

École Polytechnique Fédérale de Lausanne

Thermal properties of MOFs from molecular dynamics

Master Thesis

Lucas De Almeida Kloss

November 13, 2026

Acknowledgments

I would like to thank professor Michele Ceriotti for giving me the opportunity to do my master thesis in his group. I had previously learned about molecular dynamics throughout some of my courses, but never had the opportunity to apply it in a project. I had also previously other courses on machine learning, but again never had the opportunity to apply it in a real project. When I discovered the COSMO group on the EPFL website, I was immediately interested in the work they were doing and immediately knew I wanted to do my master thesis there. I have gained a lot of experience in numerical simulations, algorithms, machine learning and software development during this project, which will be very useful for my future career.

I would like to also thank Qianjun Xu for his help and guidance throughout the project, and for always being available to answer my questions. The discussions we had were very helpful to resolve the problems I encountered and to save time during the simulations. I would like to also thank Guillaume Fraux for explaining certain concepts when I was contributing to parts of the C++ API of the Metatomic and Metatensor libraries. I wanted to get better at programming and he gave me the opportunity to do so by letting me contribute and by reviewing my code.

AI usage disclosure

AI tools were used during this project mainly for code generation and syntax correction. The models used throughout the project were OpenAI's ChatGPT 5.6 Luna, 5.6 Terra, 5.6 Sol, 6 Astra and 6.1 Sol.

Contents

1 Introduction **3**

2 Theory **4**
- 2.1 Quantum Mechanics . . . 4
- 2.2 Density Functional Theory . . . 5
- 2.3 Molecular dynamics . . . 6
- 2.4 Heat Capacity . . . 8
- 2.5 Point Edge Transformer . . . 10

3 Methods **14**
- 3.1 MOFs . . . 14
- 3.2 Guest molecules . . . 14
- 3.3 PET-MAD and PET-SOL . . . 15
- 3.4 NPT simulation . . . 15
- 3.5 Hybrid C_P . . . 16
- 3.6 Uncertainty quantification . . . 17

4 Results **20**
- 4.1 MOF-5 . . . 20
- 4.2 MOF-303 . . . 24
- 4.3 ZIF-8 . . . 26

5 Conclusion **28**

References **29**

A Transformers and attention **33**
- A.1 Tokenization and embedding . . . 33
- A.2 Attention mechanism . . . 34
- A.3 Multi-head attention . . . 35

B Sparse automatic differentiation **36**
- B.1 The Jacobian, JVPs, and VJPs . . . 36
- B.2 Recovering multiple columns with a single evaluation . . . 37
- B.3 Graph coloring . . . 37
- B.4 Decompression . . . 38
- B.5 The Hessian and HVPs . . . 39
- B.6 ASD in PET . . . 40
- B.7 Hessian truncation . . . 41

C Temperature and pressure control **42**
- C.1 Thermostat . . . 42
- C.2 Barostat . . . 43

1 Introduction

Metal-organic frameworks (MOFs) are porous crystalline materials built with metal nodes and organic linkers, as their name suggests. These highly porous materials have a large internal surface area and can adsorb a wide variety of molecules, making them useful for applications such as gas storage. As these materials adsorb and release gas molecules, they also exchange heat with their surroundings. Therefore, studying the material's heat capacity is important for understanding these processes. The heat capacity of a material is defined as the amount of energy required to raise the temperature of the material by one degree. However, experimental data on the heat capacity of MOFs is not widely available. Especially when considering all the potential gas molecules that can be adsorbed in the pores and the wide temperature ranges.

As a result, it would be useful to develop an accurate computational method to predict the heat capacity of MOFs. But because these materials and their adsorbed molecules are large and complex systems, there are many challenges when trying to model them accurately and efficiently. For empty MOFs, the material being only in the solid phase, the heat capacity can be computed using the harmonic approximation, which is relatively easy and accurate. The problem arises when inserting guest molecules into the pores of the MOF. These molecules can move around more and interact with the framework as well as the other guest molecules. These interactions are anharmonic, which makes the harmonic approximation less accurate. To make things worse, quantum effects can also play a role in these systems, especially at low temperatures. When trying to run molecular dynamics simulations to capture all these effects, with a wide temperature range and different loadings, the computational cost can become very high.

Capturing both harmonic and anharmonic effects at an affordable computational cost is the central challenge attemped by [22]. They used and evaluated a hybrid approach for MOF-5 at different CH_4 loadings and temperatures, combining classical molecular dynamics (MD) with a harmonic correction to account for quantum effects. Our goal with this project is to first try to reproduce their results using PET [**?**], a machine-learned interatomic potential (MLIP). Once we have validated the approach, we will then extend it to other MOFs (MOF-303 and ZIF-8) and guest molecules (CO_2 and H_2O). Additionally, we will also use a method proposed by [43], that uses automatic differentiation and the sparsity of the Hessian matrix, in order to compute the harmonic heat capacity at a lower computational cost.

2 Theory

2.1 Quantum Mechanics

In order to better understand the methods used in this work, we briefly review some important quantum mechanics concepts. These are necessary to understand the Born–Oppenheimer approximation, which is what allows us to run classical molecular dynamics simulations with a potential energy surface generated. But also density functional theory, which is the method used to generate the reference energies, forces, and stresses for the machine-learned potentials.

2.1.1 Schrödinger's equation

In (non-relativistic) quantum mechanics, the fundamental equation that describes a quantum system is the Schrödinger equation [3]. Let $\hat{\mathcal{H}}$ be a Hermitian operator called Hamiltonian, and $|\Psi\rangle$ a vector in a Hilbert space called a state vector. The Schrödinger equation can be written as an eigenvalue problem for the Hamiltonian operator

$$\hat{\mathcal{H}}|\Psi_n\rangle = E_n|\Psi_n\rangle, \tag{2.1}$$

where the eigenvalues E_n are real numbers and correspond to the energy of the state $|\Psi_n\rangle$. More generally, if $|\Psi(t)\rangle$ or $\hat{\mathcal{H}}(t)$ are time-dependent, the Schrödinger equation can be written as

$$i\hbar\frac{\partial}{\partial t}|\Psi(t)\rangle = \hat{\mathcal{H}}|\Psi(t)\rangle, \tag{2.2}$$

where $\hbar$ is the reduced Planck constant and i is the imaginary unit. The position representation is a way of expressing the state of a quantum system in terms of the position. For instance, consider the Hermitian position operator $\hat{\boldsymbol{x}}$. It posesses eigenstates $\hat{\boldsymbol{x}}|\boldsymbol{x}\rangle = \boldsymbol{x}|\boldsymbol{x}\rangle$, with eigenvalues $\boldsymbol{x}$, which is a real number corresponding to the position of the particle. If we project $|\Psi(t)\rangle$ onto $|\boldsymbol{x}\rangle$, we obtain the wavefunction $\Psi(\boldsymbol{x}, t) = \langle\boldsymbol{x}|\Psi(t)\rangle$, which is a complex-valued function of position and time. In this representation, the momentum operator $\hat{\boldsymbol{p}}$ is also a Hermitian operator, and it's given by $\hat{\boldsymbol{p}} = -i\hbar\nabla_{\boldsymbol{x}}$. If we write $\hat{\mathcal{H}} = \hat{p}^2/2m + \hat{V}$ in the position representation, we have $\hat{\mathcal{H}}(\boldsymbol{x}) = -\frac{\hbar^2}{2m}\nabla^2_{\boldsymbol{x}} + V(\boldsymbol{x})$. The Schrödinger equation in the position representation is then

$$i\hbar\frac{\partial\Psi(\boldsymbol{x}, t)}{\partial t} = \left[-\frac{\hbar^2}{2m}\nabla^2_{\boldsymbol{x}} + V(\boldsymbol{x})\right]\Psi(\boldsymbol{x}, t). \tag{2.3}$$

2.1.2 The Born–Oppenheimer approximation

Let us now consider a system of N_e electrons and N_n nuclei with coordinates $\boldsymbol{r}$ and $\boldsymbol{R}$ respectively. The Hamiltonian of this system in the position representation can be written as

$$\hat{\mathcal{H}}(\boldsymbol{r}, \boldsymbol{R}) = \hat{T}_\mathrm{e} + \hat{T}_\mathrm{n} + \hat{V}_\mathrm{ee}(\boldsymbol{r}) + \hat{V}_\mathrm{en}(\boldsymbol{r}; \boldsymbol{R}) + \hat{V}_\mathrm{nn}(\boldsymbol{R}), \tag{2.4}$$

where $\hat{T}_\mathrm{e}$ and $\hat{T}_\mathrm{n}$ are the kinetic energy operators for the electrons and nuclei, and $\hat{V}_\mathrm{ee}$, $\hat{V}_\mathrm{en}$, and $\hat{V}_\mathrm{nn}$ are the electron–electron, electron–nuclear, and nuclear–nuclear interaction potentials. Because the mass of the nuclei is much larger than that of the electrons, we can decouple the electronic and nuclear motions using the Born–Oppenheimer approximation [4]. This approximation consits in separating the electronic part of the Hamiltonian $\hat{\mathcal{H}}_\mathrm{el}(\boldsymbol{R}) = \hat{T}_\mathrm{e} + \hat{V}_\mathrm{ee} + \hat{V}_\mathrm{en}(\boldsymbol{R})$ from the

nuclear part of the Hamiltonian $\hat{\mathcal{H}}_{\text{nuc}} = \hat{T}_{\text{n}} + \hat{V}_{\text{nn}}(\boldsymbol{R})$. We then solve the electronic Schrödinger equation by treating $\boldsymbol{R}$ as a fixed parameter:

$$\left[-\sum_a \frac{\hbar^2}{2m_{\text{e}}}\nabla^2_{\boldsymbol{r}_a} + V_{\text{ee}}(\boldsymbol{r}) + V_{\text{en}}(\boldsymbol{r};\boldsymbol{R})\right]\psi_0(\boldsymbol{r};\boldsymbol{R}) = \varepsilon_0(\boldsymbol{R})\psi_0(\boldsymbol{r};\boldsymbol{R}). \tag{2.5}$$

By solving this equation at $\boldsymbol{R}$, we obtain the electronic ground-state wavefunction $\psi_0(\boldsymbol{r};\boldsymbol{R})$ and its ground-state energy $\varepsilon_0(\boldsymbol{R})$. We can then use $\varepsilon_0(\boldsymbol{R})$ in the potential of the nuclear system, which is given by $U(\boldsymbol{R}) = \varepsilon_0(\boldsymbol{R}) + V_{\text{nn}}(\boldsymbol{R})$. The nuclear Schrödinger equation is then given by

$$\left[-\sum_i \frac{\hbar^2}{2M_i}\nabla^2_{\boldsymbol{R}_i} + U(\boldsymbol{R})\right]\chi_\nu(\boldsymbol{R}) = E_\nu\chi_\nu(\boldsymbol{R}). \tag{2.6}$$

When solving this equation, we obtain the nuclear wavefunction $\chi_\nu(\boldsymbol{R})$ and its energy E_ν. The total wavefunction of the system can then be approximated as a product of the electronic and nuclear wavefunctions

$$\Psi_\nu(\boldsymbol{r},\boldsymbol{R}) \approx \psi_0(\boldsymbol{r};\boldsymbol{R})\chi_\nu(\boldsymbol{R}). \tag{2.7}$$

2.2 Density Functional Theory

Within the Born–Oppenheimer approximation, the electronic ground-state energy $\varepsilon_0(\boldsymbol{R})$ must be determined for fixed nuclear coordinates $\boldsymbol{R}$. Solving Eq. (2.5) directly is computationally demanding because the many-electron wavefunction depends on all electronic coordinates. Density functional theory (DFT) instead uses the electron density $n(\boldsymbol{r})$, a function of only three spatial coordinates [8, 9].

2.2.1 The Hohenberg–Kohn theorems

The Hohenberg–Kohn theorems establish that the ground-state density determines the external potential, up to an additive constant, and hence the electronic Hamiltonian. They also provide a variational principle: the ground-state density minimizes the electronic energy functional at fixed electron number [8]:

$$\varepsilon_0(\boldsymbol{R}) = \min_{n \text{ admissible}} \left\{F[n] + \int v_{\text{ext}}(\boldsymbol{r};\boldsymbol{R})n(\boldsymbol{r})\,\mathrm{d}\boldsymbol{r}\right\}, \qquad \int n(\boldsymbol{r})\,\mathrm{d}\boldsymbol{r} = N_{\text{e}}. \tag{2.8}$$

Here v_{ext} is the electron–nuclear potential introduced in Section 2.1.2, and the universal functional $F[n]$ contains the electronic kinetic and electron–electron interaction energies. The theorems establish this density formulation but do not give a practical expression for $F[n]$.

2.2.2 The Kohn–Sham equations

The HK theorems do not provide a method for finding the ground-state density. The Kohn–Sham (KS) approach introduces a system of noninteracting electrons with the same $n(\boldsymbol{r})$ as the interacting system [9]

$$n(\boldsymbol{r}) = \sum_a f_a |\phi_a(\boldsymbol{r})|^2, \qquad \sum_a f_a = N_\mathrm{e}, \tag{2.9}$$

where $\phi_a(\boldsymbol{r})$ are the KS orbitals and f_a are their occupation numbers. The universal functional is written as $F[n] = T_\mathrm{s}[n] + E_\mathrm{H}[n] + E_\mathrm{xc}[n]$, where T_s is the noninteracting kinetic energy, E_H is the classical electrostatic energy of the electron density, and E_xc contains the remaining interaction and kinetic-energy contributions. Minimizing the energy with respect to the orbitals gives

$$\left[-\frac{\hbar^2}{2m_\mathrm{e}}\nabla^2 + v_\mathrm{ext}(\boldsymbol{r};\boldsymbol{R}) + v_\mathrm{H}[n](\boldsymbol{r}) + v_\mathrm{xc}[n](\boldsymbol{r})\right]\phi_a(\boldsymbol{r}) = \epsilon_a \phi_a(\boldsymbol{r}), \tag{2.10}$$

with $v_\mathrm{H} = \delta E_\mathrm{H}/\delta n$ and $v_\mathrm{xc} = \delta E_\mathrm{xc}/\delta n$. Because these potentials depend on the density, the equations are solved self-consistently by alternating orbital and density updates until convergence. The challenge to having good DFT results is to find a good approximation for the exchange-correlation functional $E_\mathrm{xc}[n]$.

2.3 Molecular dynamics

The potential $U(\boldsymbol{R}^N)$ appearing in the nuclear Schrödinger equation (2.6) is the Born-Oppenheimer potential energy surface. To know how the system evolves over time, we should solve this equation. However, because nuclei are much heavier than electrons, they can be treated as classical particles. Therefore, we can model their dynamics using classical mechanics, which is the basis of molecular dynamics (MD) simulations. Let $\boldsymbol{R}^N = (\boldsymbol{R}_1, \ldots, \boldsymbol{R}_N)$ and $\boldsymbol{P}^N = (\boldsymbol{P}_1, \ldots, \boldsymbol{P}_N)$ be the positions and momenta of the N nuclei in the system. The classical Hamiltonian is then given by

$$\mathcal{H}(\boldsymbol{R}^N, \boldsymbol{P}^N) = \sum_{i=1}^{N} \frac{|\boldsymbol{P}_i|^2}{2M_i} + U(\boldsymbol{R}^N), \tag{2.11}$$

Once we know the electronic ground-state energy $\varepsilon_0(\boldsymbol{R})$, the dynamics of the nuclei can be treated classically. Using Hamilton's equations, we can derive the equations of motion for the nuclei:

$$\dot{\boldsymbol{R}}_i = \frac{\partial \mathcal{H}}{\partial \boldsymbol{P}_i} = \frac{\boldsymbol{P}_i}{M_i}, \qquad \dot{\boldsymbol{P}}_i = -\frac{\partial \mathcal{H}}{\partial \boldsymbol{R}_i} = -\nabla_{\boldsymbol{R}_i} U(\boldsymbol{R}^N). \tag{2.12}$$

Since analytically solving a system of $6N$ coupled ordinary differential equations is impossible, we need to use numerical methods to integrate these equations of motion in order to find the trajectories of the nuclei. One of the most popular methods is the Runge–Kutta method, which is a family of iterative methods for solving ordinary differential equations. However, in molecular dynamics, the most commonly used method is the velocity Verlet algorithm [1]. The velocity Verlet algorithm is a symplectic integrator, which means that it preserves the geometric structure of the phase space and conserves energy over long timescales. By writing the velocities as $\boldsymbol{V}_i(t) = \dot{\boldsymbol{R}}_i(t)$ and accelarations as $\boldsymbol{A}_i(t) = \dot{\boldsymbol{P}}_i(t)/M_i$, we can update the positions of the particles using

$$\boldsymbol{R}_i(t + \Delta t) = \boldsymbol{R}_i(t) + \boldsymbol{V}_i(t)\Delta t + \frac{1}{2}\boldsymbol{A}_i(t)\Delta t^2, \tag{2.13}$$

Using this new position, we can calculate $\boldsymbol{A}_i(t + \Delta t)$ and update the velocities using

$$\boldsymbol{V}_i(t+\Delta t) = \boldsymbol{V}_i(t) + \frac{1}{2}\left[\boldsymbol{A}_i(t) + \boldsymbol{A}_i(t+\Delta t)\right]\Delta t. \tag{2.14}$$

If we want to simulate a system with constant N, V and E, we can use the velocity Verlet algorithm to integrate the equations of motion. However, in this project we want to simulate a system with constant N, P and T, which aims to reproduce the conditions of a laboratory experiment. We have therefore to make use of statistical mechanics to sample the appropriate ensemble.

2.3.1 Statistical Mechanics

The goal of statistical mechanics is to relate the microscopic properties of a system to its macroscopic behavior. In the context of molecular dynamics, we are interested in the properties of the NPT ensemble, which is characterized by constant number of particles N, pressure P_{ext}, and temperature T. The quantity of interest in this case is the enthalpy, which is defined as $H = E + PV$, where E is the internal energy, P is the pressure, and V is the volume. The internal energy E being the Hamiltonian $\mathcal{H}(\boldsymbol{R}^N, \boldsymbol{P}^N)$, the enthalpy is then given by

$$H(\boldsymbol{R}^N, \boldsymbol{P}^N) = \sum_{i=1}^{N} \frac{|\boldsymbol{P}_i|^2}{2m_i} + U(\boldsymbol{R}^N) + PV. \tag{2.15}$$

In the NPT ensemble, both internal energy and enthalpy fluctuate as the system exchanges heat and work with its surroundings. In this ensemble, a macroscopic observable A is given by the ensemble average of its microscopic observable $A(\boldsymbol{R}^N, \boldsymbol{P}^N)$ as

$$\langle A \rangle = \int_0^\infty \mathrm{d}V \int_{\text{cell}(V)} \mathrm{d}^{3N}R \int \mathrm{d}^{3N}P A(\boldsymbol{R}^N, \boldsymbol{P}^N) e^{-\beta[\mathcal{H}(\boldsymbol{R}^N, \boldsymbol{P}^N) + PV]}, \tag{2.16}$$

The particles are confined to a periodic simulation cell of volume V and shape $\boldsymbol{h}$. Additionally, we assume the ergodic hypothesis, which states that the infinite time average of an observable along a trajectory is equal to its ensemble average. So the integral (2.16) can be replaced by a time average along a sufficiently long trajectory

$$\langle A \rangle = \lim_{\tau \to \infty} \frac{1}{\tau} \int_0^\tau A\big(\boldsymbol{R}^N(t), \boldsymbol{P}^N(t)\big)\, \mathrm{d}t. \tag{2.17}$$

In practice, we can't run the simulation for ever, so we need to run the simulation for long enough to ensure that the system has reached equilibrium and explored the relevant phase space. We can then replace the integral by a finite sum over the trajectory frames, giving the time average

$$\langle A \rangle \approx \frac{1}{M} \sum_{j=1}^{M} A\big(\boldsymbol{R}^N(t_j), \boldsymbol{P}^N(t_j)\big), \tag{2.18}$$

The temperature and the pressure are controlled by a thermostat and a barostat to sample the NPT ensemble. The corresponding equations and their implementations are described in Appendix C.

2.4 Heat Capacity

The heat capacity of a material measures the energy required to raise its temperature by 1 K. In the context of MOFs, the heat capacity is an important property to understand their thermal behavior and stability. The heat capacity can be defined at constant volume C_V or at constant pressure C_P. In this project, we are interested in the heat capacity at constant pressure C_P, since it is what is measured in experiments and is relevant for practical applications. There are different ways to compute it depending on the level of theory used to describe the system. We show here the different methods used in this project.

2.4.1 C_P

As said earlier, the heat capacity is defined as the amount of energy required to raise the temperature of a system by 1 K. In the NPT ensemble, this can be expressed as the temperature derivative of the enthalpy

$$C_P(T) = \left(\frac{\partial \langle H \rangle}{\partial T}\right)_{N,P}, \qquad \langle H \rangle = \langle E + PV \rangle. \tag{2.19}$$

The average enthalpy $\langle H \rangle$ is obtained from the MD trajectory as a time average, as shown in Eq. (2.18). If we perform MD simulations at different temperatures, we can compute it by using finite differences

$$C_P(T) \simeq \frac{\langle H \rangle_{T+\Delta T} - \langle H \rangle_{T-\Delta T}}{2\Delta T}. \tag{2.20}$$

2.4.2 Dulong-Petit law

The heat capacity of a solid at constant volume can be estimated using the Dulong–Petit law. In a solid, the atoms are arranged in a lattice and can vibrate around their equilibrium positions. The potential energy of the system can be approximated as a quadratic function of the atomic displacements. The Hamiltonian is then a sum of quadratic terms in the atomic displacements and momenta

$$\mathcal{H}(\boldsymbol{R}^N, \boldsymbol{P}^N) = \sum_{i=1}^{N} \frac{|\boldsymbol{P}_i|^2}{2M_i} + \frac{1}{2}\sum_{i=1}^{N}(k_{i,x}R_{i,x}^2 + k_{i,y}R_{i,y}^2 + k_{i,z}R_{i,z}^2), \tag{2.21}$$

where $k_{i,x}$, $k_{i,y}$, and $k_{i,z}$ are the force constants for the i-th atom in the x, y, and z directions respectively. Using the equipartition theorem, we can compute the contribution of each quadratic degree of freedom to the energy. We get that the average energy is $\langle E \rangle_{\text{har}} = 3Nk_{\text{B}}T$. For constant volume, the heat capacity is then given by

$$C_{V,\text{cl}}^{\text{har}} = \left(\frac{\partial \langle E \rangle_{\text{har}}}{\partial T}\right)_{N,V} = 3Nk_{\text{B}}. \tag{2.22}$$

In other words, the classical heat capacity of a solid is constant and doesn't depend on the temperature. This law doesn't hold well at lower temperatures, where quantum effects become important.

2.4.3 Quantum harmonic heat capacity from the Hessian

The harmonic approximation can also be used to compute the quantum harmonic heat capacity of a solid. Let's consider N atoms displaced by $\boldsymbol{u}$ from their equilibrium position $\boldsymbol{R}_0$, the harmonic approximation gives

$$U(\boldsymbol{R}_0 + \boldsymbol{u}) \simeq U_0 + \frac{1}{2}\boldsymbol{u}^\mathsf{T}\boldsymbol{H}\boldsymbol{u}, \qquad H_{i\alpha,j\beta} = \left.\frac{\partial^2 U}{\partial R_{i\alpha}\partial R_{j\beta}}\right|_{\boldsymbol{R}_0}. \tag{2.23}$$

Using the diagonal mass matrix $\boldsymbol{M}$, which contains the masses of the atoms, we can define the mass-weighted Hessian $\boldsymbol{D}$ and its eigenvectors $\boldsymbol{e}_s$ as

$$\boldsymbol{D} = \boldsymbol{M}^{-1/2}\boldsymbol{H}\boldsymbol{M}^{-1/2}, \qquad \boldsymbol{D}\boldsymbol{e}_s = \omega_s^2\boldsymbol{e}_s. \tag{2.24}$$

By finding the eigenvalues ω_s^2 of $\boldsymbol{D}$, we can write a quantum version of the Hamiltonian (2.21) in terms of the normal mode coordinates $\hat{Q}_s$ and momenta $\hat{P}_s$

$$\hat{\mathcal{H}}_{\text{har}} = U_0 + \sum_s \left(\frac{\hat{P}_s^2}{2} + \frac{1}{2}\omega_s^2\hat{Q}_s^2\right). \tag{2.25}$$

Because this equation is of the form of a quantum harmonic oscillator and each mode is independent, we can solve the Schrödinger equation (2.1) for each mode independently. Their energy levels are given by $E_{s,n} = \hbar\omega_s(n + 1/2)$, where n is a non-negative integer. Using the canonical partition function

$$Z = e^{-\beta U_0}\prod_s\sum_{n=0}^{\infty} e^{-\beta\hbar\omega_s(n+1/2)} = e^{-\beta U_0}\prod_s \frac{e^{-\beta\hbar\omega_s/2}}{1 - e^{-\beta\hbar\omega_s}}, \tag{2.26}$$

we can compute the average energy of the system as

$$\langle E\rangle_{\text{har}} = -\frac{\partial \ln Z}{\partial\beta} = U_0 + \sum_s \hbar\omega_s\left(\frac{1}{2} + \frac{1}{e^{\beta\hbar\omega_s} - 1}\right). \tag{2.27}$$

The heat capacity at constant volume is then given by

$$C_{V,\text{qn}}^{\text{har}}(T) = \frac{\partial\langle E\rangle_{\text{har}}}{\partial T} = \sum_s k_\text{B}x_s^2\frac{e^{x_s}}{(e^{x_s} - 1)^2}, \qquad x_s = \frac{\hbar\omega_s}{k_\text{B}T}. \tag{2.28}$$

Therefore, by computing the Hessian of the atomic system and its eigenvalues, we can simply compute the quantum harmonic heat capacity of the system. As $T \to \infty$, the heat capacity approaches the Dulong–Petit limit $C_{V,\text{qn}}^{\text{har}} \to 3Nk_\text{B}$.

2.5 Point Edge Transformer

A machine-learned interatomic potential (MLIP) approximates the energy of a system of atoms based on their positions and types, making it a relatively inexpensive and generalizable approach. This section describes the Point Edge Transformer (PET) architecture [32], which is used to generate the machine-learned potentials PET-MAD and PET-SOL.

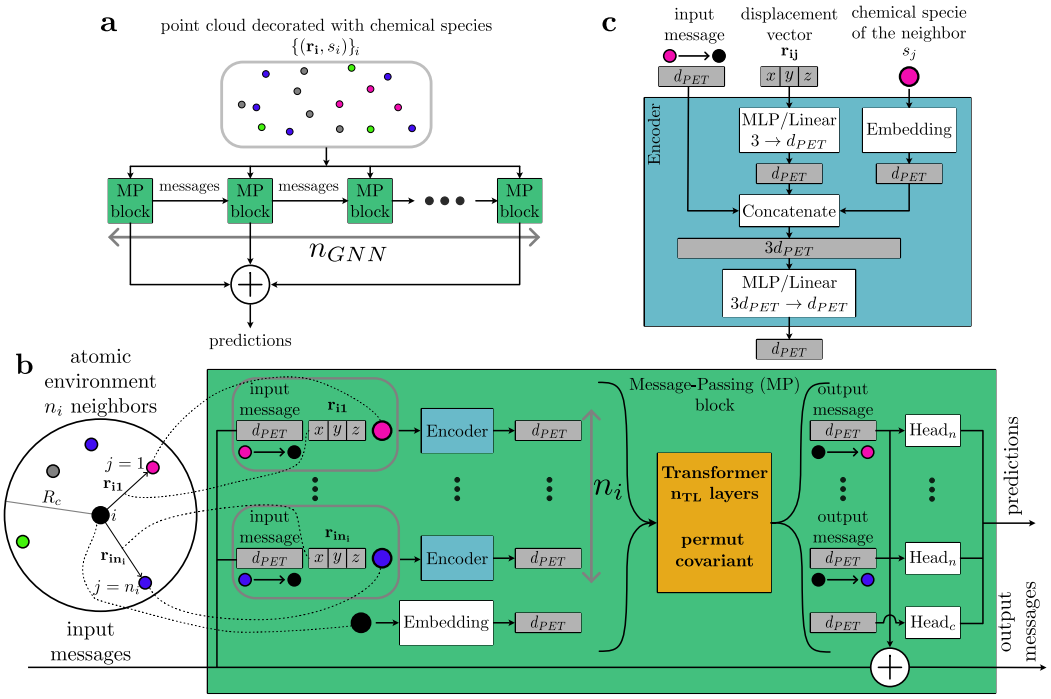


Figure 1: Point Edge Transformer architecture, reproduced from Ref. [32]. Each message-passing block embeds geometry, species, and incoming messages, processes the local tokens with a transformer, and produces both target contributions and outgoing directed messages. Local transformer depth n_{TL} and inter-environment message-passing depth n_{GNN} play different roles.

Model overview The first step in the PET architecture is to define a local environment $A_i = \{(\mathbf{r}_{ij}, s_j) | r_{ij} < R_c\}$ for an atom i, where $\mathbf{r}_{ij}$ represents the position of atom j relative to atom i, s_j is the species of atom j, which is a categorical variable that could correspond to the the atomic number Z_j, and R_c is the cutoff radius, which limits the number of neighbors considered in the environment. The environment includes all neighboring atoms j within this cutoff distance from the central atom i. The same local block is applied to all atomic environments, usually batched for parallel evaluation, and the total energy sums their learned contributions. Once the environment is defined, the PET architecture processes it through n_{GNN} message-passing (MP) blocks as shown in Figure 1 a. The first MP block takes as input only the atomic environment A_i and produces an output (here the energy), as well as an output message. The next MP block takes as input the same atomic environment A_i as well as the output message from the previous block. It will then output a new energy and a new output message. This process is reapeated for all n_{GNN} MP blocks, with each block taking as input the same atomic environment A_i and the output message from the previous block. The final output of the PET architecture is the sum of all energy contributions from each MP block.

Encoding and embedding At each message-passing block k, PET converts the species, geometry, and incoming messages of the environment A_i into tokens of width d_{PET} (Figure 1b–c). Following Appendix A of Ref. [32], the central and neighbor species use separate learned embedding maps, $E_c^{(k)}(s_i), E_n^{(k)}(s_j) \in \mathbb{R}^{d_{\mathrm{PET}}}$. These maps allow the same chemical species to have different representations in the central and neighbor roles. The relative position of neighbor j is embedded through a learned linear layer followed by a SiLU activation:

$$E_r^{(k)}(\boldsymbol{r}_{ij}) = \sigma\big(L^{(k)}(\boldsymbol{r}_{ij})\big)\,, \qquad L^{(k)} : \mathbb{R}^3 \longrightarrow \mathbb{R}^{d_{\mathrm{PET}}}. \tag{2.29}$$

This maps the three Cartesian components to the same feature width as the species embeddings. The resulting features encode spatial geometry and are not required to transform as physical vectors under rotations. For blocks after the first, the neighbor also supplies the incoming message $\tilde{x}_{ij}^{(k-1)} \in \mathbb{R}^{d_{\mathrm{PET}}}$, sent from j to i in the preceding block. Equation (5) of Ref. [32] combines this message with the species and position embeddings to form the neighbor token:

$$t_j^{(k)} = M^{(k)}\Big[\mathrm{concatenate}\Big(\tilde{x}_{ij}^{(k-1)}, E_n^{(k)}(s_j), E_r^{(k)}(\boldsymbol{r}_{ij})\Big)\Big]\,. \tag{2.30}$$

Concatenation produces $3d_{\mathrm{PET}}$ components. The one-hidden-layer MLP $M^{(k)}$ mixes these inputs and compresses them to a token of width d_{PET}. The central atom i is implicit in the notation $t_j^{(k)}$. In the first block, no incoming messages exist, so the message input is omitted and $M^{(1)}$ maps $2d_{\mathrm{PET}}$ features to d_{PET}. The central token is simply $E_c^{(k)}(s_i)$ and has no displacement input. An environment containing m neighbors therefore supplies $m+1$ tokens to the transformer: one for each neighbor and one for the center.

Transformer layers and attention The tokens constructed above are processed by n_{TL} transformer layers. As explained in Appendix A, self-attention lets each token use information from every other token in the environment. Consequently, a neighbor's representation can depend on the surrounding geometry as well as its own displacement from the center.

A neighbor must lose its influence smoothly as it approaches R_c. Otherwise, removing its token at the cutoff would abruptly change the attention weights and their normalization, potentially making the predicted energy discontinuous. PET therefore introduces a smooth cutoff, given in Appendix F.1, Eq. (9), of Ref. [32]:

$$f_c(r \mid R_c, \Delta_{R_c}) = \begin{cases} 1, & r \le R_c - \Delta_{R_c}, \\ \frac{1}{2}\left[\tanh\left(\frac{1}{x+1} + \frac{1}{x-1}\right) + 1\right], & R_c - \Delta_{R_c} < r < R_c, \\ 0, & r \ge R_c, \end{cases}$$
$$x = 2\frac{r - R_c + 0.5\Delta_{R_c}}{\Delta_{R_c}}. \tag{2.31}$$

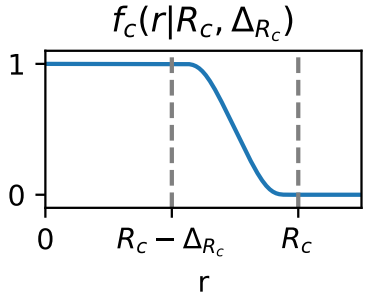


Figure 2: Smooth cutoff function $f_c(r \mid R_c, \Delta_{R_c})$ from Eq. (2.31) with $R_c = 5.0$ Å and $\Delta_{R_c} = 0.5$ Å. The function is unity for $r < R_c - \Delta_{R_c}$, smoothly decreases to zero over the switching interval, and is zero for $r > R_c$.

The width $\Delta_{R_c} > 0$ controls the interval over which a neighbor is switched off. The constant

branches are used at the endpoints, where the fractions in the middle branch are undefined. The function joins smoothly to these branches, although a narrow switching interval can still produce large derivatives within it.

Section 6 of Ref. [32] applies this cutoff to the ordinary softmax attention weights and then renormalizes them. Let a, b, c index tokens in the environment centered on atom i, with $c_b = f_c(r_{ib} \mid R_c, \Delta_{R_c})$ for a neighbor and $c_b = 1$ for the central token. For one attention head, the modified weights and output are

$$\alpha_{ab}^{\text{cut}} = \frac{\alpha_{ab} c_b}{\sum_c \alpha_{ac} c_c}, \qquad \boldsymbol{y}_a = \sum_b \alpha_{ab}^{\text{cut}} \boldsymbol{v}_b, \tag{2.32}$$

where α_{ab} are the ordinary weights from Eq. (A.2) and $\boldsymbol{v}_b$ is the value vector. The superscript cut distinguishes the modified weights from the original ones. As a neighbor leaves, it disappears smoothly from both the weighted sum and its normalization. The central token retains unit cutoff weight, keeping the denominator positive in exact arithmetic.

Each transformer layer combines multi-head attention with token-wise feedforward operations, residual connections, and layer normalization; Figure 13 illustrates the general arrangement. The output tokens of one layer become the inputs to the next, adding processing within the same environment. Shared operations preserve permutation equivariance: reordering the neighbor tokens produces the same reordering of their outputs, so their association with individual neighbors is retained (Figure 1b).

Transformer outputs The n_{TL} transformer layers produce $n_{\text{i+1}}$ outputs, $x_i^{(k)}$ for the central token and $x_{ji}^{(k)}$ for each neighbor j. The $x_{ji}^{(k)}$ correspond to the output messages that will be fed into the next MP block. The energy prediction of the block is computed using $H_n^{(k)}(x_{ji}^{(k)})$ for the neighbor contributions and $H_c^{(k)}(x_i^{(k)})$ for the central contribution, which are learned readouts that output scalar values. Using the cutoff function $f_c(r_{ij} \mid R_c, \Delta_{R_c})$ to smoothly remove contributions from neighbors beyond the cutoff distance, the neighbor contributions are summed to produce the total energy contribution from the block:

$$y_i^{\text{bond}\,(k)} = \sum_{j \in A_i} H_n^{(k)}(x_{ji}^{(k)}) f_c(r_{ij} \mid R_c, \Delta_{R_c}). \tag{2.33}$$

The cutoff is required here as well as inside attention. Attention removes a neighbor's influence on surviving tokens, but its own output need not vanish; the readout cutoff removes its direct contribution to the energy. The alternative in the paper's Eq. (6) is

$$y_i^{\text{bond}\,(k)} = \frac{\sum_{j \in A_i} H_n^{(k)}(x_{ji}^{(k)}) f_c(r_{ij} \mid R_c, \Delta_{R_c})}{\sum_{j \in A_i} f_c(r_{ij} \mid R_c, \Delta_{R_c})}. \tag{2.34}$$

This weighted average changes sensitivity to coordination and requires special handling when all neighbor weights vanish. The sum naturally returns zero for an empty neighbor list. The two readouts are alternatives that must be selected consistently in training and inference.

Appendix B adds fitted atomic self-contributions to the network output. Writing $E_{\text{self}} = \sum_i b_{s_i}$ for that composition-dependent baseline, the following expression assembles the paper's central, edge, and baseline terms:

$$\tilde{E} = E_{\text{self}} + \sum_{k=1}^{n_{\text{GNN}}} \sum_{i} \left[H_c^{(k)}(x_i^{(k)}) + y_i^{\text{bond}\,(k)} \right]. \tag{2.35}$$

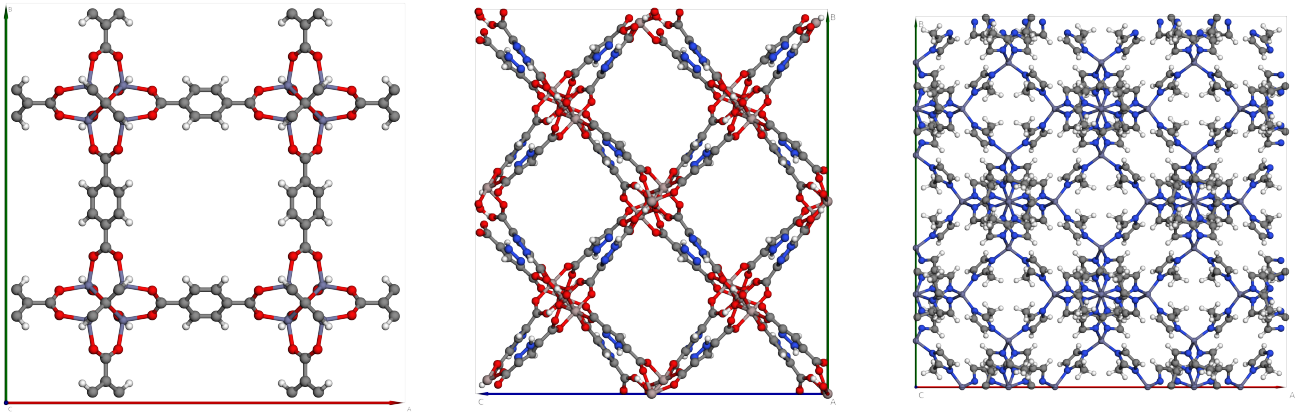
Equation (2.35) is bookkeeping assembled from the text, not a separately numbered equation in Ref. [32]; E_{self} and b_s are explanatory symbols. The total $\tilde{E}$ is the learned potential U_θ used elsewhere in this report. Atomic self-contributions are fitted by linear regression on species counts before network training, subtracted from the targets during training, and restored at inference. They have zero Cartesian derivatives at fixed composition.

All blocks contribute to the readout. Central and edge terms are learned parts of the total energy, not uniquely defined physical atomic or bond energies. An edge output depends on its environment and incoming messages, so it represents a many-body contribution; no extra pair-counting factor of one half is implied. Forces follow by differentiating the total energy through all environments, messages, attention normalizations, and cutoffs, as in Eq. (2.12).

3 Methods

3.1 MOFs

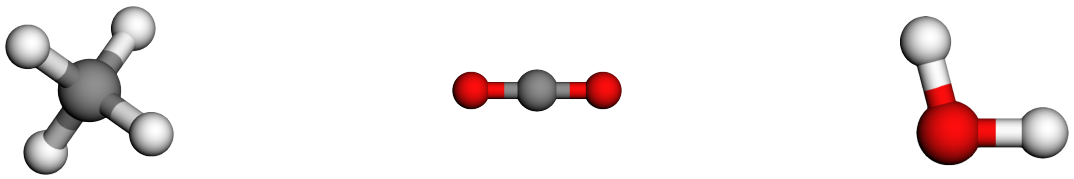
The methane comparison uses the conventional cubic MOF-5 cell, with initial lattice parameter 25.866 Å and composition $C_{192}H_{96}O_{104}Zn_{32}$ (424 host atoms). Its Zn_4O nodes are connected by benzene-1,4-dicarboxylate linkers. Adding 50, 100, or 150 methane molecules gives 674, 924, or 1174 atoms, respectively. Periodic boundary conditions apply in all three directions. MOF-303, Al(OH)(PZDC), contains aluminium–oxygen rods connected by 1H-pyrazole-3,5-dicarboxylate linkers [10]; its $2 \times 2 \times 2$ host supercell contains 1024 atoms. ZIF-8, $Zn(mIm)_2$, contains zinc centres joined by 2-methylimidazolate linkers in a sodalite-type network [11]; its $2 \times 2 \times 2$ host supercell contains 2208 atoms.



(a) MOF-5 (b) MOF-303 (c) ZIF-8

Figure 3: Empty MOfs used in this project and rendered with Chemiscope [45]. Atoms are colored by species: C (gray), H (white), O (red), Zn (blue). MOF-5 has formula $C_{192}H_{96}O_{104}Zn_{32}$, MOF-303 has formula $Al_{64}C_{128}H_{64}N_{64}O_{128}$, and ZIF-8 has formula $C_{864}H_{864}N_{576}Zn_{288}$. MOF-5 contains 424 atoms; the MOF-303 and ZIF-8 $2 \times 2 \times 2$ supercells contain 1024 and 2208 atoms, respectively.

3.2 Guest molecules

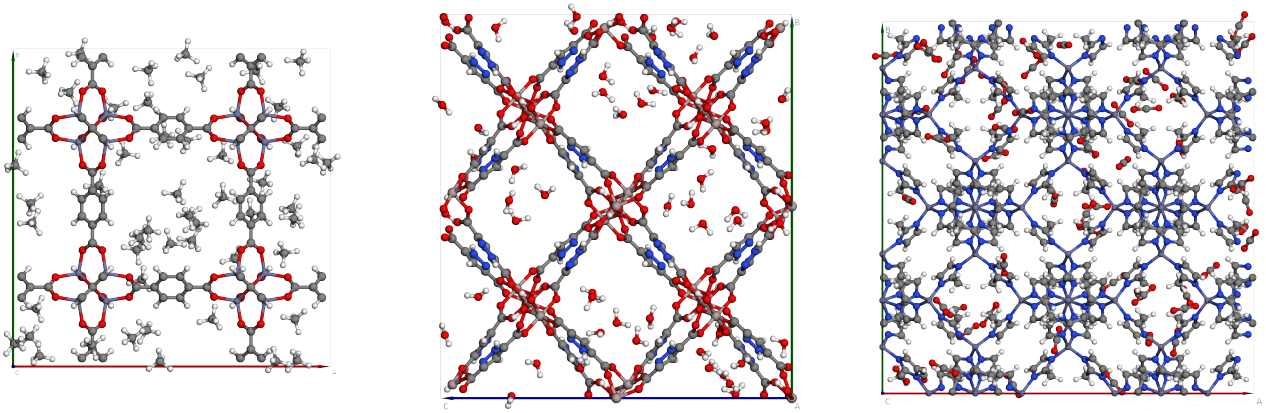


(a) CH_4 (b) CO_2 (c) H_2O

Figure 4: Guest molecules used in this project and rendered with Chemiscope [45]. H_2O is the only polar molecule, while CH_4 and CO_2 are non-polar.

The guest molecules are inserted sequentially at random positions and orientations within the MOF cell. Their placements are accepted only when all guest atoms remain at least 1.5 Å from the framework and previously inserted molecules. If the insertion becomes blocked, the previously placed guests are partially or fully removed, we then reinsert the remaining guests,

and continue until the desired number of molecules is reached. The resulting structure may be further relaxed before starting the MD simulation if initial forces are too large. Figure 5 shows representative starting configurations with guests inside the same periodic host cells as in Figure 3, before relaxation.



(a) MOF-5 + 50 CH_4 (b) MOF-303 + 50 H_2O (c) ZIF-8 + 54 CO_2

Figure 5: Guest-loaded MOFs rendered with Chemiscope [45], using the same viewing directions and element colors as the empty frameworks. MOF-5 contains 50 methane molecules (674 atoms in total), the MOF-303 $2 \times 2 \times 2$ supercell contains 50 water molecules (1174 atoms), and the ZIF-8 $2 \times 2 \times 2$ supercell contains 54 carbon dioxide molecules (2370 atoms). These are insertion configurations before relaxation.

After successfully inserting the guest molecules, we need to assign the initial velocities to all the atoms in the system. This is done by sampling from a Maxwell-Boltzmann distribution at the target temperature

$$P(\boldsymbol{v}_i) = \left(\frac{m_i}{2\pi k_\mathrm{B} T}\right)^{3/2} \exp\left(-\frac{m_i |\boldsymbol{v}_i|^2}{2 k_\mathrm{B} T}\right). \tag{3.1}$$

3.3 PET-MAD and PET-SOL

Two pretrained PET models are used in this work, PET-MAD and PET-SOL. They are both based on the same PET architecture, as described in Section 2.5, but they were trained on different datasets. More precisely, PET-MAD-1.5 was trained on the MAD-1.5 dataset [30], spanning 102 elements and including molecules, clusters, crystals, and low-dimensional structures. Its reference energies, forces, and periodic stresses were computed with all-electron r^{2}SCAN DFT. PET-SOL on the other hand was trained on a dataset of 17 solvents, 78 cations and 35 anions. The models are available in different sizes, but here we will use PET-MAD-1.5-s and PET-SOL-s, which are smaller than the full model and will make the MD simulations faster. Both models have the same number of parameters and share common settings summarized in the following table.

3.4 NPT simulation

As said earlier, we will be running NPT simulations to compute the heat capacity of the loaded MOF. They will be carried out using the LAMMPS package [2]. In particular, the LAMMPS implementation used is the one developed in [44], which allows us to run MD simulations using

Setting	Value
R_c, Δ_{R_c}	8.0 Å, 0.5 Å
n_{GNN}, n_{TL}	3, 1
d_{PET}	256
Attention heads	8
Transformer feed-forward width	512
Node feature width, `d_node`	1024

Table 1: Architecture settings of PET-MAD and PET-SOL

MLIPs such as PET-MAD and PET-SOL. The following table summarizes the main parameters used in the NPT simulations.

Parameter	Value
T, ΔT	175–425 K, 25 K
P	1 bar
τ_T, τ_P	100, 1000 fs
t, Δt	500 ps, 0.5 fs
t_{eq}	100 ps
Frame interval	0.05 ps

Table 2: Parameters used in all NPT simulations

The simulations will run for 500 ps with a timestep of 0.5 fs, and the first 100 ps will be discarded to allow for equilibration. This is equivalent to 1 million timesteps, with the first 200,000 discarded. Because this would require a lot of storage, we will only save the trajectory every 0.05 ps, which is equivalent to every 100 timesteps, giving a total of 8000 frames per trajectory. Our results will be reported for temperatures between 200 K and 400 K, but because we are using finite differences in (2.20), we will also need to run simulations at 175 K and 425 K, which will only be used for the finite difference calculations. Each temperature MD simulation will be run on a NVIDIA Tesla V100 GPU, which has 32 GB of VRAM.

3.5 Hybrid C_P

An empty MOF is a rigid structure, so it is appropriate to use the harmonic approximation to compute its heat capacity. For lower temperatures, the heat capacity from the Dulong-Petit law (2.22) is not accurate. So in the general case, we would use the quantum harmonic heat capacity (2.28), which includes quantum effects and is more accurate at lower temperatures. Doing so gives results that are pretty close to experimental results, as we will see later. However, when including guest molecules inside the MOF, because these molecules are loose and their interactions with the framework and with each other are not harmonic, the harmonic approximation is not valid anymore. Because the loaded MOF has different contributions to the heat capacity from the framework and from the guest molecules, Kapil et al. [22] proposed a hybrid approach to compute the heat capacity of the loaded MOF. The hybrid approach uses the classical anharmonic heat capacity from MD simulations and adds a harmonic correction

$$C_P^{\text{hyb}}(T) \simeq C_{P,\text{cl}}^{\text{anh}}(T) + C_{V,\text{qn}}^{\text{har}}(T) - C_{V,\text{cl}}^{\text{har}}(T). \tag{3.2}$$

The first term $C_{P,\mathrm{cl}}^{\mathrm{anh}}(T)$ is computed using (2.20) from the MD simulations. The second term $C_{V,\mathrm{qn}}^{\mathrm{har}}(T)$ is computed using the quantum harmonic heat capacity from (2.28). The third term $C_{V,\mathrm{cl}}^{\mathrm{har}}(T)$ is computed using the Dulong-Petit law from (2.22). This construction retains classical anharmonicity while approximating the quantum correction harmonically. It neglects quantum anharmonic effects and assumes that the quantum change in $C_P - C_V$ is small.

3.5.1 Computing the frequencies

To compute the harmonic heat capacity in (3.2), we need to get the frequencies from the eigenvalues of the Hessian matrix. To do so, we select the last frame of the MD trajectory at the highest temperature of 400 K and relax the structure using a FIRE optimizer. Once the structure is relaxed and the forces are minimized, we can compute the eigenvalues of the Hessian matrix. The frequencies used in (2.28) will be the same for all temperatures, so only the temperature included in the formula will change. The efficiently compute the Hessian matrix, we use sparse automatic differentiation, which is a method that uses the sparsity of the Hessian to compute it efficiently. In the appendix B, the principles of the sparse automatic differentiation are explained in more detail with a simple example. The actual implementation is done using the methods from [43], and the asdex library [**?**] is used to compute the Hessian matrix. The eigenvalues of the Hessian matrix can be negative, in which case they are discarded. As explained in B, to speed up the calculations, we can truncate the Hessian by selecting a number of hops $k < K$, where K is the number of hops needed to recover the full Hessian. In this project with PET-S, we will use a number of hops $k = 3$, which shows a very small relative error compared to the full Hessian with $K = 7$, when computing the heat capacity [43].

3.6 Uncertainty quantification

When using a machine learning model to predict quantities, there are many sources of uncertainty. The model may not have been trained on data that is representative of the system being studied. It may not capture all relevant physics, which depends on the first principles methods used when doing DFT calculations. If we wanted to have better uncertainty estimates, we would need to have the model trained on different datasets and use these models to predict the same quantity. We could then have a better idea on the uncertainty of the model predictions. Additionally, we could also build complete different models with different architectures and train them with different datasets. Combining the results from each model to have a final prediction and uncertainty estimate is what we call ensemble learning. In traditional machine learning, ensemble learning is a technique that combines the predictions of multiple models to improve overall performance and provide more robust uncertainty estimates. In practice, ensemble learning is widely used in various machine learning tasks, including regression, classification, and forecasting. One of the most efficient models are tree-based models, such as Random Forests and Gradient Boosting Machines, which combine the predictions of multiple decision trees to improve accuracy and reduce overfitting.

3.6.1 Last-layer prediction rigidity

Training multiple models, especially in the case of MLIPs, is an expensive task. Instead of training multiple models, we can use the last-layer prediction rigidity (LLPR) [39] method to estimate

the uncertainty of a single model. The idea is that instead of training multiple models, we take the model we are using for the simulation and we create new models by just changing its last layer. The models with the different last layers could be considered as different models that were trained on different datasets. This method is cheap and efficient, and it has been shown to provide good uncertainty estimates for MLIPs [39]. For central energy $\bar{U} = \boldsymbol{w}_0^{\mathrm{T}} \boldsymbol{f}(\boldsymbol{x})$ and covariance-set features $\boldsymbol{F}$,

$$\begin{aligned} \sigma^2_{\mathrm{LLPR}}(\boldsymbol{x}) = \alpha^2 \boldsymbol{f}(\boldsymbol{x})^{\mathrm{T}} \\ \times \left(\boldsymbol{F}^{\mathrm{T}}\boldsymbol{F} + \lambda \boldsymbol{I}\right)^{-1} \boldsymbol{f}(\boldsymbol{x}). \end{aligned} \tag{3.3}$$

The regularizer λ stabilizes weakly constrained directions; α is calibrated against energy residuals on a disjoint reference-labeled set using the base potential's electronic-structure definition. Predictions cannot calibrate their own error. The exported `energy_uncertainty` is an energy standard deviation. Each new model is generated by sampling the last-layer weight vector $\boldsymbol{w}_i$ from the following distribution

$$\boldsymbol{w}_i \sim \mathcal{N}\left[\boldsymbol{w}_0, \alpha^2 \left(\boldsymbol{F}^{\mathrm{T}}\boldsymbol{F} + \lambda \boldsymbol{I}\right)^{-1}\right], \tag{3.4}$$

The model is then created by multiplying the sampled weight vector with the weights from all the previous layers of the model, which are kept frozen.

$$U^{(i)}(\boldsymbol{x}) = \boldsymbol{w}_i^{\mathrm{T}} \boldsymbol{f}(\boldsymbol{x}) \tag{3.5}$$

The advantage of this method is that it is not only cheap to generate the models, but also we do not need to rerun the simulations with them. The model used to run the simulation can be considered to be the central model and the mean of all the other models generated by sampling the last layer is equal to the central model. That means that we can use the same trajectory generated by the central model to evaluate the quantities with the other models, which will be treated as uncertainties.

3.6.2 Cumulative expansion approximation

We can then compute the difference between the potential of the members and the central model as well as the enthalpy of each member.

$$\Delta U_t^{(i)}(T) = U_t^{(i)}(T) - \bar{U}_t(T), \tag{3.6}$$

$$H_t^{(i)}(T) = K_t + U_t^{(i)}(T) + PV_t, \tag{3.7}$$

In particular, if we want to calculate the ensemble average of the enthalpy as a function of time, we can use the following equation

$$\left\langle H^{(i)}(T) \right\rangle = \frac{\sum_t e^{-\beta \Delta U_t^{(i)}(T)} H_t^{(i)}}{\sum_t e^{-\beta \Delta U_t^{(i)}(T)}}, \tag{3.8}$$

It is done using the Boltzmann factors of the energy differences between the central model and the members. When the energy differences $\Delta U_t^{(i)}$ are small, the direct reweighting can be approximated by a first-order cumulant expansion approximation (CEA):

$$\begin{aligned}\left\langle H^{(i)}(T)\right\rangle &\approx \left\langle H^{(i)}(T)\right\rangle_{\bar{U}} - \beta\,\mathrm{cov}_{\bar{U}}\left(H^{(i)}(T), \Delta U^{(i)}(T)\right), \\ \beta &= (k_{\mathrm{B}}T)^{-1}.\end{aligned} \tag{3.9}$$

The CEA is a useful approximation when the energy differences are small, but it is not exact. Direct reweighting remains a diagnostic: $N_{\mathrm{eff}} = 1/\sum_t \widetilde{w}_t^2$ should remain adequate, and $\mathrm{var}_t(\beta\Delta U_t^{(i)}) \ll 1$ is needed for controlled first-order CEA. Applying Eq. (2.20) to each member's CEA mean enthalpies gives $C_{P,\mathrm{cl}}^{(i)}(T)$. We can then simply calculate the uncertainty at each temperature as the standard deviation of the heat capacities of all the members:

$$\sigma_{C_{P,\mathrm{cl,anh}}}(T) = \left\{\frac{1}{M-1}\sum_{i=1}^{M}\left[C_{P,\mathrm{cl}}^{(i)}(T) - \overline{C}_{P,\mathrm{cl}}(T)\right]^2\right\}^{1/2}. \tag{3.10}$$

This term corresponds to the classical anharmonic heat capacity uncertainty. To compute the uncertainty of the quantum harmonic heat capacity, we need first to select the last frame in the simulation, relax the structure and then compute the Hessian. We then calculate the eigenvalues of the Hessian matrix, which correspond to the vibrational frequencies of the system. We will compute the Hessian for each model and its eigenvalues, and then we will compute the quantum harmonic heat capacity for each model using Eq. (2.28). Finally, we can compute the uncertainty of the quantum harmonic heat capacity as the standard deviation of the heat capacities of all the models:

$$\sigma_{C_{V,\mathrm{qn,har}}}(T) = \left\{\frac{1}{M-1}\sum_{i=1}^{M}\left[C_{V,\mathrm{har}}^{(i)}(T) - \overline{C}_{V,\mathrm{har}}(T)\right]^2\right\}^{1/2}. \tag{3.11}$$

The total uncertainty of the hybrid heat capacity is then obtained by combining the uncertainties of the classical anharmonic and quantum harmonic heat capacities in quadrature:

$$\sigma_{C_P}(T) = \sqrt{\sigma^2_{C_{P,\mathrm{cl,anh}}}(T) + \sigma^2_{C_{V,\mathrm{qn,har}}}(T)}. \tag{3.12}$$

4 Results

4.1 MOF-5

4.1.1 CH_4

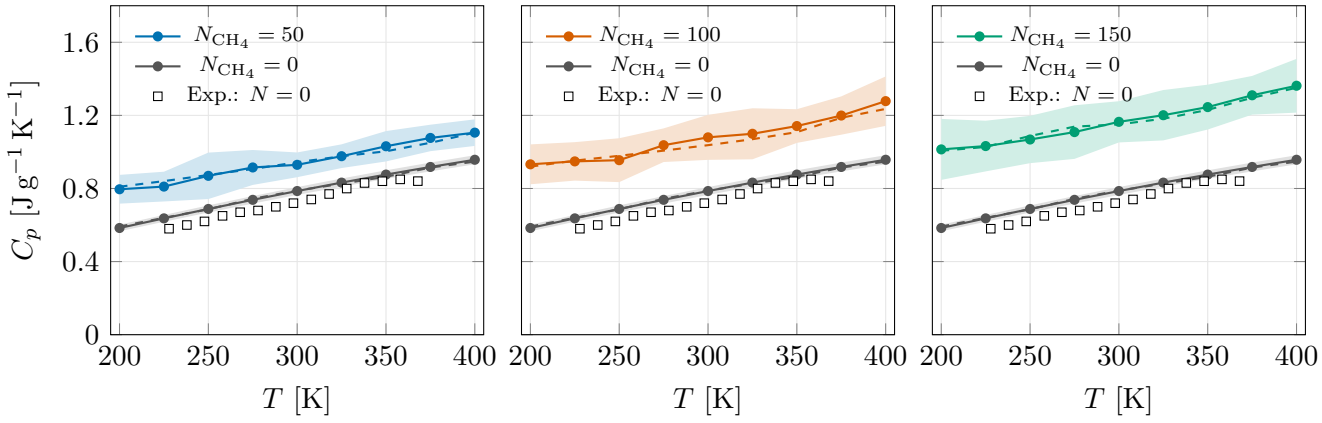


Figure 6: MOF-5 C_P with 50, 100, and 150 CH_4 molecules per cell using PET-MAD. Dashed lines show the results of Kapil *et al.* [22]. Open squares show experimental empty-MOF-5 data from Ming *et al.* [24]. Shading denotes standard uncertainty; details are given in Section 3.6.

Kapil *et al.* [22] used QuickFF, a potential specifically designed for MOFs and trained on DFT data. PET-MAD gives very similar heat capacities, with differences of at most 0.043 J/g/K at 300 K. This agreement shows that a general-purpose potential can reproduce the results of a specialized MOF potential. If PET-MAD is faster and more efficient than QuickFF, its broader applicability could make it a better choice for future simulations of other systems, particularly larger MOFs or larger supercells.

The agreement is especially striking because the reference calculations used PIMD, which is significantly more expensive than classical MD. Combining PET-MAD with classical MD and the hybrid method reproduces the QuickFF/PIMD heat capacities closely, providing accurate results at a much lower computational cost.

A major difference from the work of Kapil *et al.* concerns the uncertainty calculation: the present estimates include both MD and model uncertainty. Figure 6 shows the uncertainty bands obtained using 64 committee members. The three loadings are presented in separate panels to avoid overlapping bands.

The empty framework has much narrower uncertainty bands than the loaded systems. Its heat capacity requires no MD simulation, so the uncertainty is given only by Equation 3.12 and is much smaller than the MD uncertainty. The bands also widen with increasing loading. This may reflect the stable configuration of the empty framework, for which different committee members give similar energies. In the loaded systems, the guest molecules are not in a stable configuration, and changes to the model's last layer affect the guest–guest and guest–host interaction energies more strongly, increasing the model uncertainty.

PET-SOL also reproduces the QuickFF results accurately (Figure 7). At 300 K, the differences are at most 0.043 J/g/K. Despite its specific training on solvents, PET-SOL offers no advantage over PET-MAD, which was trained on a more general dataset. Its overall error is slightly larger.

The agreement of both potentials with QuickFF supports extending classical MD and the

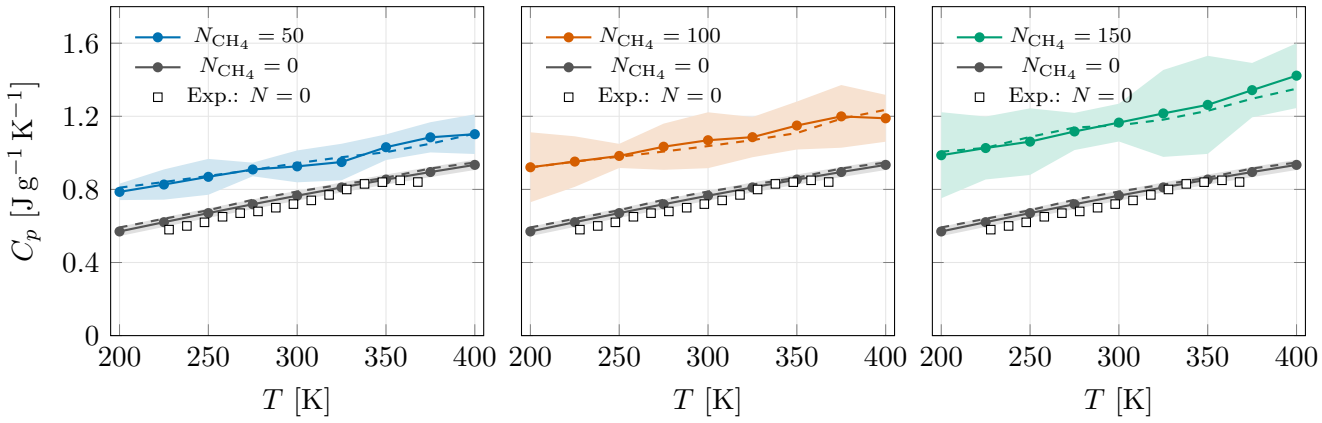


Figure 7: MOF-5 C_P with 50, 100, and 150 CH_4 molecules per cell using PET-SOL. Dashed lines show the results of Kapil *et al.* [22]. Open squares show experimental empty-MOF-5 data from Ming *et al.* [24]. Shading denotes standard uncertainty; details are given in Section 3.6.

hybrid method to other MOFs and guest molecules. To conserve computational resources, the remaining simulations use only PET-MAD because of its broader applicability and smaller errors. Table 3 summarizes the deviations from the reference data.

Model	CH_4/cell	MAE	RMSE	Bias	max $\lvert e_i \rvert$
PET-MAD	50	0.0140	0.0178	0.0000	0.0300
PET-MAD	100	0.0256	0.0287	0.0194	0.0428
PET-MAD	150	0.0150	0.0165	0.0041	0.0299
PET-SOL	50	0.0171	0.0204	-0.0027	0.0348
PET-SOL	100	0.0201	0.0258	0.0097	0.0464
PET-SOL	150	0.0297	0.0353	0.0147	0.0714

Table 3: Deviations of the MOF-5/CH_4 hybrid heat capacities from the digitized QuickFF/PIMD reference of Kapil *et al.* [22]. Each row uses nine simulation temperatures from 200 to 400 K in 25 K steps; reference values are linearly interpolated to these temperatures without extrapolation. With $e_i = c_{P,i}^{\mathrm{sim}} - c_{P,i}^{\mathrm{ref}}$, the mean absolute error (MAE) is $\langle \lvert e_i \rvert \rangle$, the root mean square error (RMSE) is $\sqrt{\langle e_i^2 \rangle}$, the bias (mean error) is $\langle e_i \rangle$, and max $\lvert e_i \rvert$ is the largest absolute deviation. Positive bias indicates overestimation. All deviations are in J/g/K and compare central estimates without uncertainty weighting.

4.1.2 CO_2 and H_2O

No experimental heat-capacity data are available for MOF-5 loaded with CO_2, and Kapil *et al.* [22] did not study this system. The calculated curve behaves qualitatively like the methane curve, but appears to approach a plateau and is very close to the empty-framework heat capacity at 400 K. If this trend continued at higher temperatures, the empty-framework heat capacity would eventually exceed that of the loaded system, which would be physically incorrect. The high-temperature CO_2 results therefore require caution.

Kapil *et al.* [22] averaged multiple MD simulations at each temperature to obtain a more accurate heat-capacity curve. Time constraints limited this project to one simulation per temperature, leaving room for improvement.

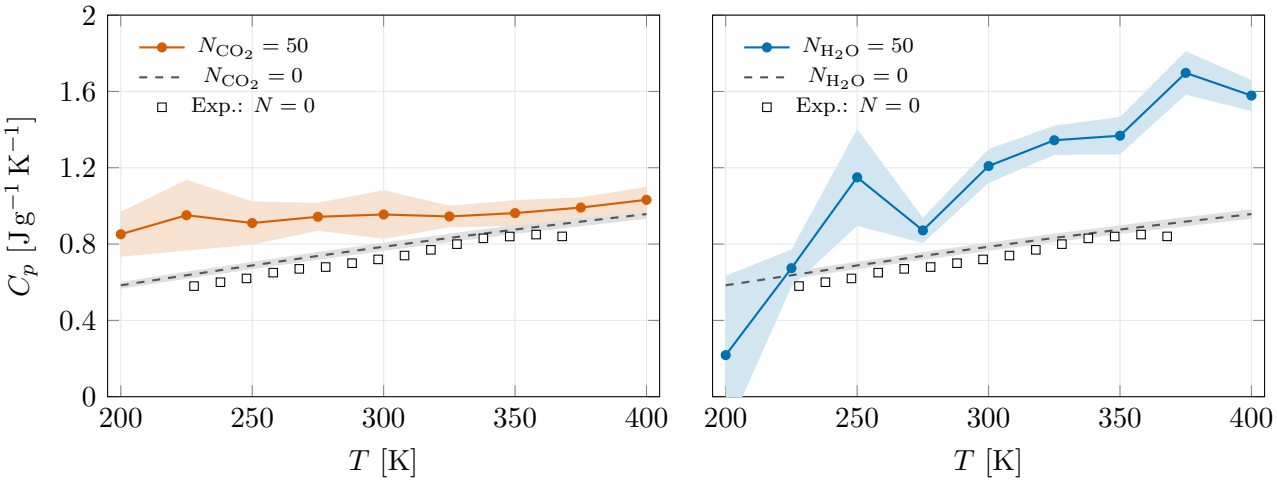


Figure 8: MOF-5 heat capacities with 50 CO_2 molecules (left) and 50 H_2O molecules (right) per cell using PET-MAD. Each panel shows the total hybrid c_P and the empty-framework harmonic c_V reference. Loaded curves use the host-plus-guest mass; empty references use the host mass. Open squares show experimental empty-MOF-5 data from Ming *et al.* [24]. Shading denotes standard uncertainty; details are given in Section 3.6. Both y-axes start at zero; uncertainty bands extending below zero are clipped at the axis boundary.

For MOF-5 with 50 H_2O molecules, the heat capacity at 200 K lies below the empty-framework reference. Since heat capacity is expected to increase with the number of molecules, this unexpected result could indicate a problem with the low-temperature simulations.

Water is expected to span three phases—solid, liquid, and gas—over 200–400 K. A solid–liquid transition is expected between 250 and 275 K, where the calculated heat capacity decreases substantially rather than increasing with temperature as usually expected. This decrease could indicate an enthalpy discontinuity associated with a phase transition. A smaller discontinuity between 350 and 375 K could similarly correspond to the expected liquid–gas transition. However, phase transitions are notoriously difficult to capture with MD, and further analysis is needed to determine whether PET can reproduce them. Experimental comparison would help assess these results, but few data are available for water in MOF-5.

PET-SOL could be an alternative to PET-MAD for this system because water is polar and acts as a solvent, and PET-SOL was trained specifically on solvents. Time constraints prevented this comparison, but testing it could establish whether it gives better results. The same discrepancies are absent from the curves for the nonpolar guests CH_4 and CO_2, suggesting that PET-MAD may be less suitable for MOFs containing polar molecules, particularly at low temperatures.

Figure 9 presents the trajectory diagnostics at 200 K. Although the mean temperature is 200.10 K, the mean enthalpy decreases by 0.559 eV per cell from the 100–300 ps interval to the 300–500 ps interval, while the mean volume increases by 0.187%. The system therefore continues to relax during the measurement interval despite maintaining the target temperature.

Recomputing the hybrid heat capacity over the early and late intervals with the same harmonic quantum correction gives 0.363 J/g/K and 0.073 J/g/K, respectively, compared with 0.218 J/g/K over the full interval. The derivative at 200 K depends on the mean enthalpies at 175 and 225 K, so this sensitivity concerns the equilibration of both the neighboring-temperature trajectories and the 200 K trajectory itself.

In nine frames sampled from the 200 K trajectory, O–H distances range from 0.915 to 1.091 Å.

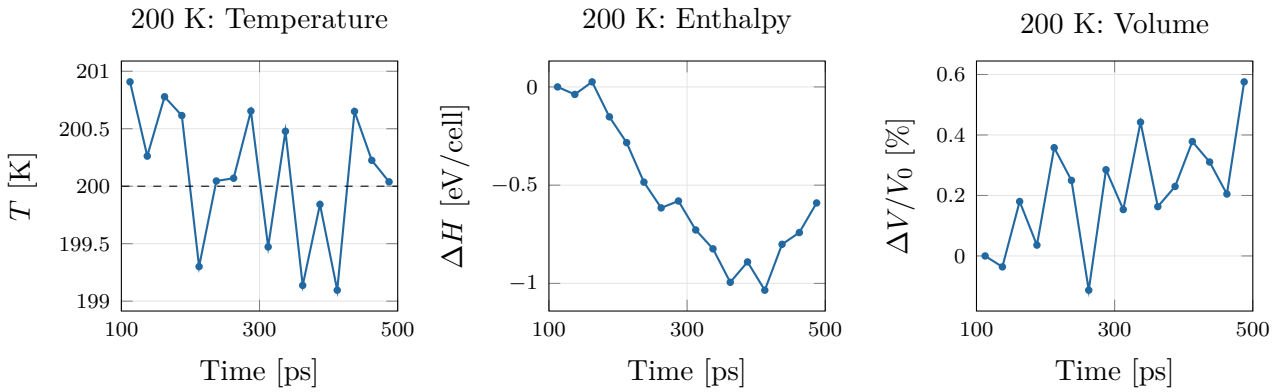


Figure 9: Trajectory diagnostics for MOF-5 with 50 H_2O molecules at 200 K. Points are nonoverlapping 25 ps block means over the 100–500 ps measurement interval. The dashed line marks the target temperature. Enthalpy and volume changes are measured relative to their first block means; V_0 is the first block mean volume. The blocks summarize one trajectory and are not independent replicas.

Each water hydrogen remains closest to its original oxygen, even when framework oxygens are included in the periodic search. These spot checks show no obvious dissociation in the sampled frames but cannot exclude events between samples. The enthalpy drift and sensitivity to the averaging window demonstrate a convergence limitation; they neither establish a phase transition nor fully explain the low heat capacity.

4.2 MOF-303

4.2.1 CH_4, CO_2, and H_2O

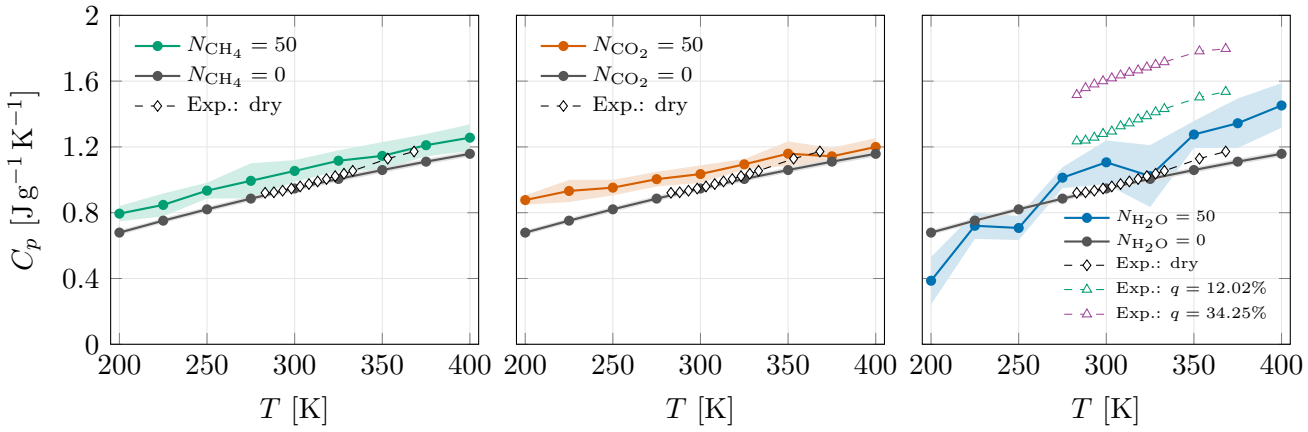


Figure 10: PET-MAD heat capacities for MOF-303 with 50 CH_4 (left), 50 CO_2 (middle), and 50 H_2O (right) molecules per $2 \times 2 \times 2$ supercell. Each panel shows the total hybrid c_P and empty-framework harmonic c_V reference. Loaded curves use the host-plus-guest mass; empty references use the host mass. The methane panel uses its campaign-specific empty Hessian. Shading denotes standard uncertainty. Dashed experimental curves connect dry-MOF-303 measurements in each panel and water loadings $q = 12.02\%$ and 34.25% in the right panel [23]; experimental uncertainties were not supplied. Here q is the fraction of saturated water loading, and the simulation corresponds to $q = 16.76\%$. Dry measurements refer to the empty framework. All y-axes start at zero.

Abu El-maaty *et al.* [23] report water loading as a fraction of the saturated loading, $q = m_{\text{water}}/m_{\text{water,sat}}$. Their dry and saturated sample masses of 7.3990 and 10.5358 mg give a saturated water-to-dry-framework mass ratio of 0.42395. In the simulation, 50 water molecules in the 1024-atom host give a water-to-host mass ratio of 0.07106, corresponding to $q = 0.16761$. This composition lies between the experimental samples with $q = 0.1202$ and $q = 0.3425$. The simulated heat capacity would therefore be expected to lie between their curves, but instead falls below both. Reproducing the measurements accurately is difficult because the exact water loading in the experimental MOF-303 sample is uncertain, whereas the simulation contains a fixed, known number of guest molecules.

As in MOF-5, the water-loaded heat capacity falls below the empty-framework reference at low temperatures. The curve is not strictly monotonic, and its low-temperature behavior is physically unreasonable. These results therefore require caution, particularly at low temperatures. PIMD could improve their accuracy despite its higher computational cost. An alternative potential, particularly PET-SOL, might also give better results.

Figure 11 shows the water-trajectory diagnostics at 200, 225, and 250 K, where the hybrid heat capacity lies below the empty-framework curve. The mean temperatures are 199.96, 224.91, and 249.84 K, respectively. Comparing the 300–500 ps interval with the 100–300 ps interval gives mean enthalpy changes of -0.553, -0.363, and -0.795 eV per supercell, accompanied by mean volume increases of 0.087%, 0.153%, and 0.318%. All three trajectories therefore continue to relax during the measurement interval despite maintaining their target temperatures.

Keeping the harmonic quantum correction fixed while using the early and late averaging intervals gives early/late hybrid heat capacities of $(0.358, 0.415)$, $(0.738, 0.703)$, and $(0.736, 0.679)$

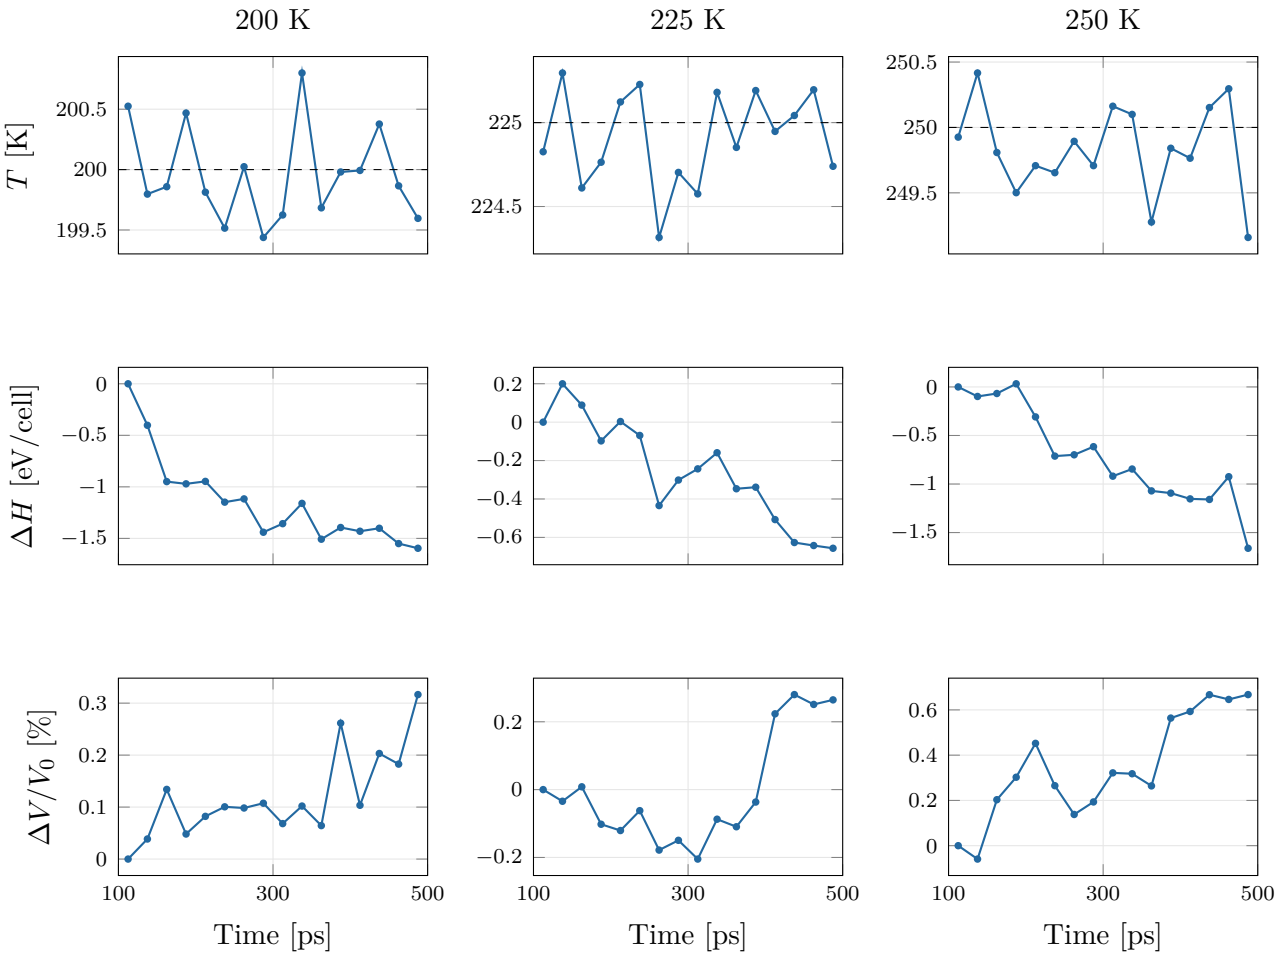


Figure 11: Trajectory diagnostics for MOF-303 with 50 H_2O molecules per $2 \times 2 \times 2$ supercell at 200, 225, and 250 K (left to right). Rows show temperature, enthalpy change, and relative volume change (top to bottom). Points are nonoverlapping 25 ps block means over 100–500 ps; dashed lines mark the target temperatures. Each trajectory's enthalpy and volume changes are referenced to its first block means. These diagnostics use one trajectory per temperature.

in J/g/K at 200, 225, and 250 K, respectively. The full-interval values are 0.386, 0.721, and 0.707 J/g/K. These estimates are correlated and do not constitute independent uncertainty bounds; each temperature derivative also depends on neighboring-temperature runs.

Structural spot checks of nine frames per trajectory give O–H distances between 0.905 and 1.067 Å, with each water hydrogen closest to its original oxygen. The diagnostics show that stationarity has not been reached, but do not demonstrate molecular breakup or fully explain the discrepancy with the empty-framework heat capacity.

4.3 ZIF-8

For CO_2-loaded ZIF-8 at 300 K, the classical contribution $c_{\mathrm{cl}} = 2.426$ and harmonic correction $\Delta c_{\mathrm{har}} = -1.460$ give $c_{\mathrm{hyb}} = 0.967 \pm 0.151\,\mathrm{J\,g^{-1}\,K^{-1}}$. The hybrid curve rises from 0.805 at 200 K to 1.200 J/g/K at 400 K (Figure 12). Comparisons with the other frameworks must account for differences in guest concentration and mass normalization.

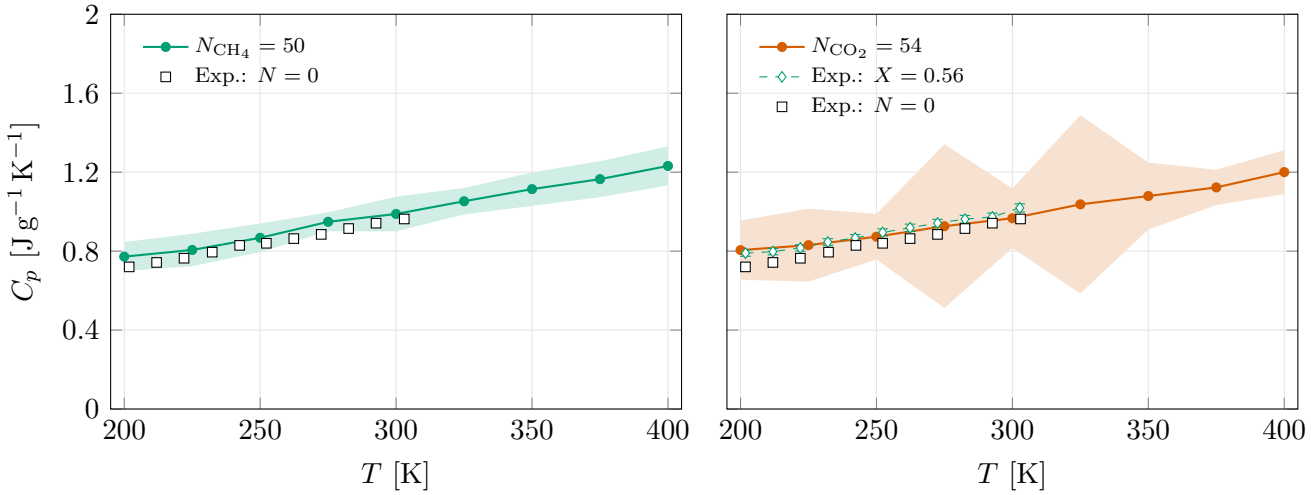


Figure 12: PET-MAD total hybrid heat capacities with standard uncertainty for 50 CH_4 (left) and 54 CO_2 (right) molecules in the ZIF-8 $2 \times 2 \times 2$ supercell. The CO_2 panel includes experimental measurements at $X = 0.56$ mol CO_2/mol ZIF-8 [26]; the simulation has $X = 0.5625$. The experimental molar values from Supporting Information Table 3 are divided by the loaded sample molar mass, 252.3 g/mol, to give the same total-sample mass basis as the simulation. Experimental error bars show 2% standard uncertainty in this temperature range; dashed lines connect measured points. The experiment reports sealed-cell C, whereas the simulation estimates C_P. Open squares show independent experimental empty-ZIF-8 (SOD) data [25], converted from molar units using 227.576 g/mol. The new digitized empty-framework datasets do not include uncertainty values, so no error bars are inferred for them. Both y-axes start at zero.

The simulated heat capacities agree surprisingly closely with the experimental measurements. On the common 200–300 K grid, linear interpolation between neighboring experimental temperatures gives signed simulation differences ranging from approximately +2.2% to −4.0%. At 300 K, the interpolated experimental value is approximately 1.007 J/g/K, compared with the simulated $0.967 \pm 0.151\,\mathrm{J\,g^{-1}\,K^{-1}}$. The 2% experimental standard uncertainty corresponds to approximately 0.020 J/g/K at this temperature and does not include a separate interpolation uncertainty.

The measurement conditions qualify this numerical agreement. The experiment contains a fixed total amount of CO_2 in a sealed pan. Neglecting pan expansion, the authors describe the

measured heat capacity as approximately constant-volume for the entire system, including any free CO_2 outside the crystal [26]. The periodic NPT simulation contains all 54 CO_2 molecules inside the framework. Moreover, the initial experimental sealing pressure is not maintained throughout the temperature scan. The 200–400 K simulation curve also samples only the upper tail of the reported 70–220 K structural anomaly.

5 Conclusion

We saw that the PET architecture is able to reproduce the results of a specialized potential for MOFs like QuickFF. As a result, we can use more general purpose potentials like PET-MAD to simulate the heat capacity of MOFs with different guest molecules. For some cases, like H_2O loaded MOFs, the results are not physically reasonable at low temperatures. We may need to further validate the results with PIMD simulations or other potentials to ensure that the results are reasonable. Given that carrying out experiments for measuring the heat capacity of MOFs is complicated, using MD simulations and the hybrid method is a good alternative to obtain the heat capacity of MOFs with different guest molecules. PET-MAD and PET-SOL give room-temperature heat capacities rising from about 0.93 to 1.16 J/g/K with methane loading. Harmonic quantization removes 59–63% of the classical term; guest mobility supports retaining anharmonic MD. Additional PET-MAD estimates at 300 K are 0.955/1.209 for CO_2/H_2O in MOF-5, 1.035/1.106 for CO_2/H_2O in MOF-303, and 0.967 J/g/K for CO_2 in ZIF-8. Experimental overlays place the MOF-303/water estimate below the measured neighboring loading curves at 300 K, while the ZIF-8/CO_2 estimate is within about 4% of the matched-composition measurement. Loading, mass-normalization, and sealed-cell versus constant-pressure conditions qualify these comparisons. Water-loaded trajectories show particularly strong drift, and all loaded spectra contain imaginary modes. Stable spectra, independent trajectory/minimum sampling, and validated paired uncertainty propagation are needed before interpreting these curves as converged predictions. Of course, because we are using finite differences to compute the classical anharmonic heat capacity, and it is the biggest source of uncertainty in our results, the obvious improvement would be to use a smaller ΔT to compute the classical anharmonic heat capacity. This would require more simulations, but it would give us a better estimate of the classical anharmonic heat capacity and reduce the uncertainty in our results.

References

[1] W. C. Swope, H. C. Andersen, P. H. Berens, and K. R. Wilson, *A computer simulation method for the calculation of equilibrium constants for the formation of physical clusters of molecules: Application to small water clusters*, Journal of Chemical Physics **76**, 637–649 (1982), doi:10.1063/1.442716.

[2] A. P. Thompson *et al.*, *LAMMPS—a flexible simulation tool for particle-based materials modeling at the atomic, meso, and continuum scales*, Computer Physics Communications **271**, 108171 (2022), doi:10.1016/j.cpc.2021.108171; LAMMPS website.

[3] C. Cohen-Tannoudji, B. Diu, and F. Laloë, *Quantum Mechanics, Volume I: Basic Concepts, Tools, and Applications*, 2nd ed., Wiley-VCH, Weinheim (2020), publisher edition.

[4] M. Born and R. Oppenheimer, *Zur Quantentheorie der Molekeln*, Annalen der Physik **84**, 457–484 (1927), doi:10.1002/andp.19270840804.

[5] M. P. Allen and D. J. Tildesley, *Computer Simulation of Liquids*, 2nd ed., Oxford University Press (2017), publisher edition.

[6] M. E. Tuckerman, *Statistical Mechanics: Theory and Molecular Simulation*, 2nd ed., Oxford University Press (2023), publisher edition.

[7] R. M. Martin, *Electronic Structure: Basic Theory and Practical Methods*, Cambridge University Press (2004), doi:10.1017/CBO9780511805769.

[8] P. Hohenberg and W. Kohn, *Inhomogeneous Electron Gas*, Physical Review **136**, B864–B871 (1964), doi:10.1103/PhysRev.136.B864.

[9] W. Kohn and L. J. Sham, *Self-Consistent Equations Including Exchange and Correlation Effects*, Physical Review **140**, A1133–A1138 (1965), doi:10.1103/PhysRev.140.A1133.

[10] N. Hanikel *et al.*, *Evolution of Water Structures in Metal-Organic Frameworks for Improved Atmospheric Water Harvesting*, Science **374**, 454–459 (2021), doi:10.1126/science.abj0890.

[11] C. L. Hobday *et al.*, *Understanding the Adsorption Process in ZIF-8 Using High Pressure Crystallography and Computational Modelling*, Nature Communications **9**, 1429 (2018), doi:10.1038/s41467-018-03878-6.

[12] C. Rowley, *Periodic Boundary Conditions in 2D* (2015), Wikimedia Commons, CC BY-SA 4.0.

[13] LAMMPS documentation, *compute pressure command*, online documentation (accessed 2 October 2026).

[14] LAMMPS documentation, *fix npt command*, online documentation (accessed 2 November 2026).

[15] LAMMPS developers, *Nosé–Hoover integrator implementation*, `FixNH::compute_scalar`, source code (accessed 15 October 2026).

[16] LAMMPS documentation, *min_style fire command*, online documentation (accessed 13 October 2026).

[17] H. Flyvbjerg and H. G. Petersen, *Error Estimates on Averages of Correlated Data*, Journal of Chemical Physics **91**, 461–466 (1989), doi:10.1063/1.457480.

[18] M. Ceriotti, M. Parrinello, T. E. Markland, and D. E. Manolopoulos, *Efficient Stochastic Thermostatting of Path Integral Molecular Dynamics*, Journal of Chemical Physics **133**, 124104 (2010), doi:10.1063/1.3489925.

[19] R. Diestel, *Graph Theory*, 6th ed., Graduate Texts in Mathematics **173**, Springer (2025), author's textbook site.

[20] A. H. Gebremedhin, A. Pothen, A. Tarafdar, and A. Walther, *Efficient Computation of Sparse Hessians Using Coloring and Automatic Differentiation*, INFORMS Journal on Computing **21**, 209–223 (2009), doi:10.1287/ijoc.1080.0286.

[21] A. Hill, *Graph Coloring*, ASDEX documentation, online documentation.

[22] V. Kapil *et al.*, *Modeling the Structural and Thermal Properties of Loaded Metal–Organic Frameworks: An Interplay of Quantum and Anharmonic Fluctuations*, Journal of Chemical Theory and Computation **15**, 3237–3249 (2019), doi:10.1021/acs.jctc.8b01297.

[23] A. E. Abu El-maaty, O. Saoudi, A. S. Abdelrazik, and R. Ben-Mansour, *Experimental investigation of specific heat and thermal stability of MOF-303 at different water vapor loading and various temperatures*, International Communications in Heat and Mass Transfer **165**, 109088 (2025), doi:10.1016/j.icheatmasstransfer.2025.109088.

[24] Y. Ming, J. Purewal, D. Liu, A. Sudik, C. Xu, J. Yang, M. Veenstra, K. Rhodes, R. Soltis, J. Warner, M. Gaab, U. Müller, and D. J. Siegel, *Thermophysical properties of MOF-5 powders*, Microporous and Mesoporous Materials **185**, 235–244 (2014), doi:10.1016/j.micromeso.2013.11.015.

[25] P. F. Rosen, J. J. Calvin, M. S. Dickson, A. D. Katsenis, T. Friščić, A. Navrotsky, N. L. Ross, A. I. Kolesnikov, and B. F. Woodfield, *Heat capacity and thermodynamic functions of crystalline forms of the metal-organic framework zinc 2-methylimidazolate, Zn(MeIm)$_2$*, The Journal of Chemical Thermodynamics **136**, 160–169 (2019), doi:10.1016/j.jct.2019.05.008.

[26] P. F. Rosen, M. S. Dickson, J. J. Calvin, N. L. Ross, T. Friščić, A. Navrotsky, and B. F. Woodfield, *Thermodynamic Evidence of Structural Transformations in CO_2-Loaded Metal–Organic Framework Zn(MeIm)$_2$ from Heat Capacity Measurements*, Journal of the American Chemical Society **142**, 4833–4841 (2020), doi:10.1021/jacs.9b13883. Experimental values and uncertainties are from Supporting Information Table 3, doi:10.1021/jacs.9b13883.s001.

[27] A. Vaswani *et al.*, *Attention Is All You Need*, Advances in Neural Information Processing Systems **30** (2017), arXiv:1706.03762.

[28] T. B. Brown *et al.*, *Language Models are Few-Shot Learners*, Advances in Neural Information Processing Systems **33**, 1877–1901 (2020), arXiv:2005.14165.

[29] A. Mazitov *et al.*, *PET-MAD as a Lightweight Universal Interatomic Potential for Advanced Materials Modeling*, Nature Communications **16**, 10653 (2025), doi:10.1038/s41467-025-65662-7.

[30] C. Malosso *et al.*, *High-quality, high-information dataset for universal atomistic machine learning*, arXiv:2603.02089v1 (2026), arXiv:2603.02089v1.

[31] PET-SOL local checkpoint, `pet_sol-s-best_nostress.ckpt`, archived model and training metadata in the project `models/` directory (inspected 3 November 2026).

[32] S. N. Pozdnyakov and M. Ceriotti, *Smooth, Exact Rotational Symmetrization for Deep Learning on Point Clouds*, Advances in Neural Information Processing Systems **36** (2023), conference paper; the supplied article is arXiv:2305.19302v3 (6 February 2024), including Appendices A–H.

[33] L. Vanduyfhuys *et al.*, *QuickFF: A Program for a Quick and Easy Derivation of Force Fields for Metal–Organic Frameworks from Ab Initio Input*, Journal of Computational Chemistry **36**, 1015–1027 (2015).

[34] G. J. Martyna, D. J. Tobias, and M. L. Klein, *Constant Pressure Molecular Dynamics Algorithms*, Journal of Chemical Physics **101**, 4177–4189 (1994), doi:10.1063/1.467468.

[35] W. Shinoda, M. Shiga, and M. Mikami, *Rapid estimation of elastic constants by molecular dynamics simulation under constant stress*, Physical Review B **69**, 134103 (2004), doi:10.1103/PhysRevB.69.134103.

[36] S. Nosé, *A unified formulation of the constant temperature molecular dynamics methods*, Journal of Chemical Physics **81**, 511–519 (1984), doi:10.1063/1.447334.

[37] W. G. Hoover, *Canonical dynamics: Equilibrium phase-space distributions*, Physical Review A **31**, 1695–1697 (1985), doi:10.1103/PhysRevA.31.1695.

[38] G. J. Martyna, M. L. Klein, and M. Tuckerman, *Nosé–Hoover Chains: The Canonical Ensemble via Continuous Dynamics*, Journal of Chemical Physics **97**, 2635–2643 (1992), doi:10.1063/1.463940.

[39] F. Bigi, S. Chong, M. Ceriotti, and F. Grasselli, *A prediction rigidity formalism for low-cost uncertainties in trained neural networks*, Machine Learning: Science and Technology **5**, 045018 (2024), doi:10.1088/2632-2153/ad805f.

[40] G. Imbalzano *et al.*, *Uncertainty estimation for molecular dynamics and sampling*, Journal of Chemical Physics **154**, 074102 (2021), doi:10.1063/5.0036522.

[41] A. G. Baydin, B. A. Pearlmutter, A. A. Radul, and J. M. Siskind, *Automatic Differentiation in Machine Learning: A Survey*, Journal of Machine Learning Research **18**, 1–43 (2018), JMLR 18(153).

[42] A. Hill, G. Dalle, and A. Montoison, *An Illustrated Guide to Automatic Sparse Differentiation*, ICLR Blogposts (2025), online article.

[43] M. F. Langer, A. Hill, and M. Ceriotti, *Truncated automatic sparse differentiation for machine learning interatomic potentials*, arXiv:2609.20510 (2026), doi:10.48550/arXiv.2609.20510.

[44] F. Bigi *et al.*, *metatensor and metatomic: Foundational libraries for interoperable atomistic machine learning*, Journal of Chemical Physics **164**, 064113 (2026), doi:10.1063/5.0304911.

[45] S. Chorna, J. Lála, Q. Xu, R. K. Cersonsky, G. Fraux, and M. Ceriotti, *Chemiscope 1.0: interactive exploration of atomistic data from analysis to dissemination*, Journal of Open Source Software **11**(122), 10380 (2026), doi:10.21105/joss.10380.

A Transformers and attention

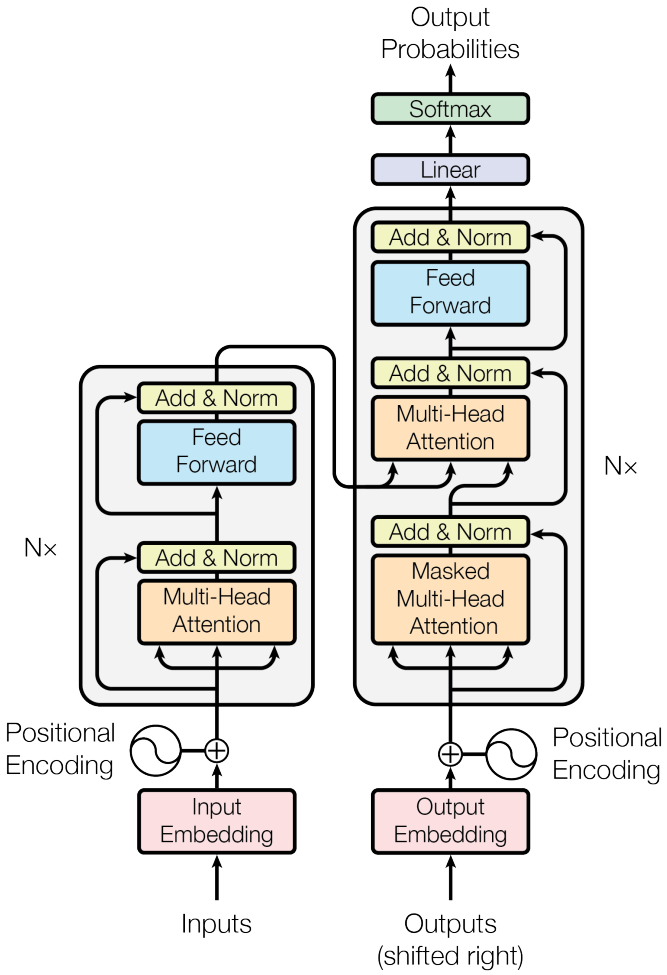


Figure 13: Transformer architecture developed in Ref. [27]. The attention mechanism is the part in gray. The part of the transformer on the left is called the encoder, and the part on the right is called the decoder. The encoder takes a sequence of tokens as input and produces a set of vectors, whose representations capture the meaning of the input tokens and their context. The decoder uses these vectors to generate an output sequence and guesses the next token in the sequence.

The transformer architecture has been revolutionary in the field of AI as a whole. It was initially developed for natural language processing (NLP) tasks, in particular for translation. But it has since been applied to a wide range of other domains, including computer vision, speech processing, code generation, protein folding. It is a state-of-the-art architecture that has been used to build some of the most powerful AI models, such as GPT-3 and GPT-4. It has also been applied to material science, in particular in MLIPs, where it has been used to predict the properties of materials and molecules.

A.1 Tokenization and embedding

The original transformer architecture [27] was designed for sentence translation. But it was then used to other NLP tasks, such as text classification and question answering. Therefore, the transformer architecture takes a text as input. However, a text is a sequence of strings, which is not numerical data and can't be directly processed by a neural network. To convert a sentence into a numerical representation, each word (or punctuation mark) in the sentence is first separated

into a token, and then each token is mapped to a vector representation, called an embedding. Each token embedding is a vector living in a high-dimensional space, and the embedding vectors are stacked together to form a matrix $\boldsymbol{X}$. The way to interpret these vectors is that each token is represented as a point in a high-dimensional space, and the distance between two points reflects the semantic similarity between the corresponding tokens. For instance, the tokens "cat" and "dog" are both animals, so their embeddings should be close to each other in the embedding space. In addition, the transformer architecture also uses positional encoding to encode the order of the tokens in the sequence. This is especially important for words that could have different meanings depending on their position in the sentence. For instance, the word "bank" could refer to a financial institution or the side of a river, and its meaning can be inferred from the surrounding words. So the embedding of the word "bank" could have different values depending on its position in the sentence, and the positional encoding helps to capture this information. These embeddings are something that has to be learned from the data, and they are usually initialized randomly and then updated during training. In the case of the PET architecture, the token embeddings are already explained in Section 2.5. So what is relevant for us is to understand the next step in the transformer architecture, which is the attention mechanism.

A.2 Attention mechanism

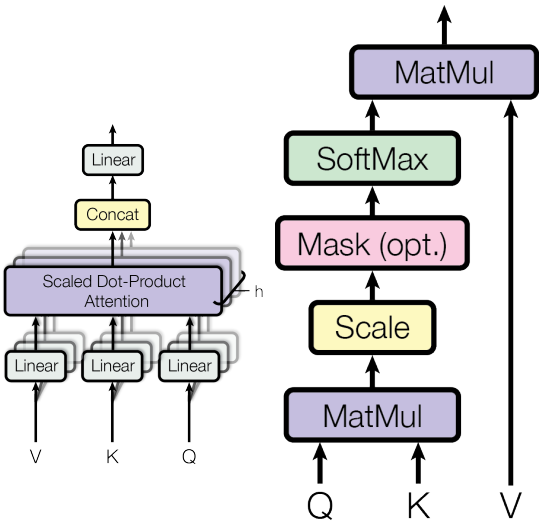


Figure 14: On the left, we have a multi-head attention mechanism, which is the orange part of the transformer architecture in Figure 13. On the right, we see the scaled dot-product attention mechanism, which is the blue part on the left figure. Pictures are taken from Ref. [27].

With our input matrix $\boldsymbol{X} \in \mathbb{R}^{n \times d}$, where n is the number of tokens and d is the embedding dimension, we compute the queries, keys, and values as follows:

$$\boldsymbol{Q} = \boldsymbol{X}\boldsymbol{W}_Q, \qquad \boldsymbol{K} = \boldsymbol{X}\boldsymbol{W}_K, \qquad \boldsymbol{V} = \boldsymbol{X}\boldsymbol{W}_V, \tag{A.1}$$

with $\boldsymbol{W}_Q, \boldsymbol{W}_K \in \mathbb{R}^{d \times d_k}$ and $\boldsymbol{W}_V \in \mathbb{R}^{d \times d_v}$. Thus $\boldsymbol{Q}, \boldsymbol{K}, \boldsymbol{V} \in \mathbb{R}^{n \times d}$. These matrices are what the attention mechanism uses to compute the attention matrix $\boldsymbol{A} \in \mathbb{R}^{n \times n}$, which contains the attention weights for each token pair.

$$\boldsymbol{A} = \mathrm{softmax}_{\mathrm{row}}\left(\frac{\boldsymbol{Q}\boldsymbol{K}^{\mathrm{T}}}{\sqrt{d_k}} + \boldsymbol{M}\right), \tag{A.2}$$

To better understand what $\boldsymbol{A}$ represents, we can look at its coefficients. Suppose we have a sentence "The animal didn't cross the street because it was too tired". The word "it" refers to "the animal" and not "the street", so the attention weight between "it" and "the animal" should be high, while the attention weight between "it" and "the street" should be low. The term $\boldsymbol{Q}\boldsymbol{K}^{\mathrm{T}}$ computes the dot product between the queries and keys, which measures the similarity between each token's query and every other token's key. The scaling factor $1/\sqrt{d_k}$ is used to prevent the dot products from growing too large. The $\boldsymbol{M}$ matrix is a mask that can be used to prevent certain tokens from attending to each other. For instance, if we are training the model to predict the next word in a sentence, we don't want the model to know the future words, otherwise it would defeat the purpose of predicting the next word, so we "mask" these words by setting their attention weights to zero. The softmax function is given by

$$\mathrm{softmax}_{\mathrm{row}}(\boldsymbol{Z})_{ab} = \frac{\exp(Z_{ab})}{\sum_{c=1}^{n}\exp(Z_{ac})}, \qquad a, b = 1, \ldots, n, \tag{A.3}$$

The softmax function normalizes the attention weights so that they sum to one across each row, which means that for each token, the attention weights represent a probability distribution over all other tokens. We add large negative values to the masked positions in $\boldsymbol{M}$ so that their softmax output is effectively zero. After computing $\boldsymbol{A}$, we can compute the output of the attention mechanism as a weighted sum of the values:

$$\boldsymbol{Y} = \boldsymbol{A}\boldsymbol{V}, \tag{A.4}$$

where $\boldsymbol{Y} \in \mathbb{R}^{n \times d_v}$ is the output of the attention mechanism. Each of the W_Q, W_K, and W_V are weight matrices that are fine-tuned during training. In the PET architecture, masking is not needed because there is no sequential ordering of the tokens, and all tokens can attend to each other.

A.3 Multi-head attention

We can extend the attention mechanism by dividing the input tokens $\boldsymbol{X}$ into subsets $\boldsymbol{X}_1, \ldots, \boldsymbol{X}_H$, such that $d_v = d/H$, and constructing attention matrices $\boldsymbol{A}_1, \ldots, \boldsymbol{A}_H$ for each subset, this is called multi-head attention. We can then concatenate the outputs of each head to form the final output:

$$\mathrm{MHA}(\boldsymbol{X}) = \mathrm{Concat}(\boldsymbol{Y}_1, \ldots, \boldsymbol{Y}_H)\boldsymbol{W}_O, \qquad \boldsymbol{W}_O \in \mathbb{R}^{Hd_v \times d}. \tag{A.5}$$

The two main advantages of using multi-head attention instead of simple attention are that it allows to make the process more parallelizable, and it allows the model to learn more complex relationships between the tokens, since it also adds more parameters to the model. A common choice is $d_k = d_v = d/H$ for $d = 256$ and eight heads this gives 32 components per head. These dimensions describe attention projections, rather than the width of a later energy readout head.

B Sparse automatic differentiation

Automatic sparse differentiation (ASD) is a method for efficiently computing derivatives of neural networks. It makes use of the sparsity of the Jacobian (or Hessian) matrix to reduce the memory usage of these matrices, by storing only the non-zero elements. The sparsity pattern of these matrices also allows us to compute the derivatives more efficiently, by avoiding unnecessary computations of zero elements. This method provides an exact computation of the derivatives, but we can make some approximations to further reduce the computations. The approximation is explained in the Hessian truncation section.

B.1 The Jacobian, JVPs, and VJPs

For a function $\boldsymbol{f} : \mathbb{R}^m \to \mathbb{R}^n$, the Jacobian matrix is defined as

$$\boldsymbol{J}(\boldsymbol{x}) = \frac{\partial \boldsymbol{f}}{\partial \boldsymbol{x}}(\boldsymbol{x}) = \begin{pmatrix} \frac{\partial f_1}{\partial x_1} & \cdots & \frac{\partial f_1}{\partial x_m} \\ \vdots & \ddots & \vdots \\ \frac{\partial f_n}{\partial x_m} & \cdots & \frac{\partial f_n}{\partial x_m} \end{pmatrix} \in \mathbb{R}^{n \times m} \tag{B.1}$$

We can compute the directional derivative of $\boldsymbol{f}$ at $\boldsymbol{x}$ along a vector $\boldsymbol{v} \in \mathbb{R}^m$ as

$$\frac{\mathrm{d}}{\mathrm{d}\epsilon} \boldsymbol{f}(\boldsymbol{x} + \epsilon \boldsymbol{v}) \bigg|_{\epsilon=0} = \boldsymbol{J}(\boldsymbol{x})\boldsymbol{v}. \tag{B.2}$$

This is what we call a *Jacobian–vector product* (JVP). It is the instantaneous rate of change of $\boldsymbol{f}$ along the direction $\boldsymbol{v}$. Since the Jacobian at a point $\boldsymbol{x}$ is a linear map, the directional derivative is linear in $\boldsymbol{v}$. This means that we can recover all the columns of the Jacobian by evaluating the JVPs along the standard basis vectors of $\mathbb{R}^m$: $\boldsymbol{e}_1, \ldots, \boldsymbol{e}_m$. Let's illustrate this with a simple example. Consider the function $\boldsymbol{f} : \mathbb{R}^4 \to \mathbb{R}^3$

$$\boldsymbol{f}(\boldsymbol{x}) = \begin{pmatrix} x_1^2 + x_2 \\ \sin(x_2) + x_3^2 \\ e^{x_4} \end{pmatrix}, \tag{B.3}$$

The Jacobian is

$$\boldsymbol{J}(\boldsymbol{x}) = \begin{pmatrix} 2x_1 & 1 & 0 & 0 \\ 0 & \cos(x_2) & 2x_3 & 0 \\ 0 & 0 & 0 & e^{x_4} \end{pmatrix}. \tag{B.4}$$

If we try evaluating the JVPs along the standard basis vectors, we get

$$\boldsymbol{J}(\boldsymbol{x})\boldsymbol{e}_1 = \begin{pmatrix} 2x_1 \\ 0 \\ 0 \end{pmatrix}, \quad \boldsymbol{J}(\boldsymbol{x})\boldsymbol{e}_2 = \begin{pmatrix} 1 \\ \cos(x_2) \\ 0 \end{pmatrix}, \quad \boldsymbol{J}(\boldsymbol{x})\boldsymbol{e}_3 = \begin{pmatrix} 0 \\ 2x_3 \\ 0 \end{pmatrix}, \quad \boldsymbol{J}(\boldsymbol{x})\boldsymbol{e}_4 = \begin{pmatrix} 0 \\ 0 \\ e^{x_4} \end{pmatrix}. \tag{B.5}$$

Which exactly recovers the columns of the Jacobian in Eq. (B.4). However, the same result could have been obtained if instead of evaluating the JVPs along the standard basis vectors of $\mathbb{R}^m$, we had used the standard basis vectors of $\mathbb{R}^n$ and evaluated the vector-Jacobian products

(VJP) along the standard basis vectors of $\mathbb{R}^n$.

$$\boldsymbol{e}_1^{\mathrm{T}}\boldsymbol{J}(\boldsymbol{x}) = \begin{pmatrix} 2x_1 & 1 & 0 & 0 \end{pmatrix}, \quad \boldsymbol{e}_2^{\mathrm{T}}\boldsymbol{J}(\boldsymbol{x}) = \begin{pmatrix} 0 & \cos(x_2) & 2x_3 & 0 \end{pmatrix}, \quad \boldsymbol{e}_3^{\mathrm{T}}\boldsymbol{J}(\boldsymbol{x}) = \begin{pmatrix} 0 & 0 & 0 & e^{x_4} \end{pmatrix}. \tag{B.6}$$

This time, we recover the rows of the Jacobian instead of the columns. The advantage of using VJPs here instead of JVPs is that if $n \ll m$, which means we can recover the Jacobian with fewer evaluations. This is particularly useful when dealing with neural networks, where the number of outputs is usually much smaller than the number of inputs. In particular, in the case MLIPs, the number of outputs is only one, the energy of the system, while the number of inputs depends on the number of atoms in the system and can be much larger. For a composition of functions $\boldsymbol{f} = \boldsymbol{h} \circ \boldsymbol{g}$, the chain rule gives the Jacobian as

$$\boldsymbol{J}_f(\boldsymbol{x})\boldsymbol{v} = \boldsymbol{J}_h(\boldsymbol{g}(\boldsymbol{x}))\big[\boldsymbol{J}_g(\boldsymbol{x})\boldsymbol{v}\big]. \tag{B.7}$$

B.2 Recovering multiple columns with a single evaluation

In some cases, we can recover multiple columns of the Jacobian with a single VJP evaluation. For instance, if we evaluate the VJP along the vector $\boldsymbol{v} = \boldsymbol{e}_1 + \boldsymbol{e}_3 + \boldsymbol{e}_4$, we get

$$\boldsymbol{J}(\boldsymbol{x})\boldsymbol{v} = \begin{pmatrix} 2x_1 \\ 2x_3 \\ e^{x_4} \end{pmatrix}. \tag{B.8}$$

The elements of this vector are the only nonzero entries in the first, third, and fourth columns of the Jacobian. By exploiting the sparsity of the Jacobian, we can recover multiple columns with a single evaluation. This is possible because column 1 has zeros in rows 2 and 3, column 3 has zeros in rows 1 and 3, and column 4 has zeros in rows 1 and 2. In other words, the nonzero entries of these columns do not overlap in any row. The nonoverlapping property is what allows us to recover multiple columns with a single evaluation. The last remaining column can be recovered with a second evaluation along $\boldsymbol{e}_2$. Thus, we can recover the entire Jacobian with only two evaluations instead of four as previously done.

B.3 Graph coloring

To better understand which columns can be grouped together, we can represent the sparsity pattern as a graph. Each column in $\boldsymbol{J}$ corresponds to a vertex in the graph and an edge connects two vertices if they have one or more nonzero entries in the same row. In our example, columns 1 and 2 have a nonzero entry in row 1, so they are connected by an edge. Same for columns 2 and 3, which have a nonzero entry in row 2. The set of vertices and edges forms a graph $G = (V, E)$, where $V = \{1, 2, 3, 4\}$ is the set of vertices (columns of $\boldsymbol{J}$) and $E = \{\{1, 2\}, \{2, 3\}\}$ is the set of edges (connections between columns). We can represent a graph with a diagram, where vertices are drawn as circles and edges as lines connecting them.

The next step is to assign a color to each vertex in the graph such that no two adjacent vertices share the same color. This process is called *graph coloring*. The goal is to find a coloring that uses the smallest number of colors possible $\chi(G)$, which is called the *chromatic number* of

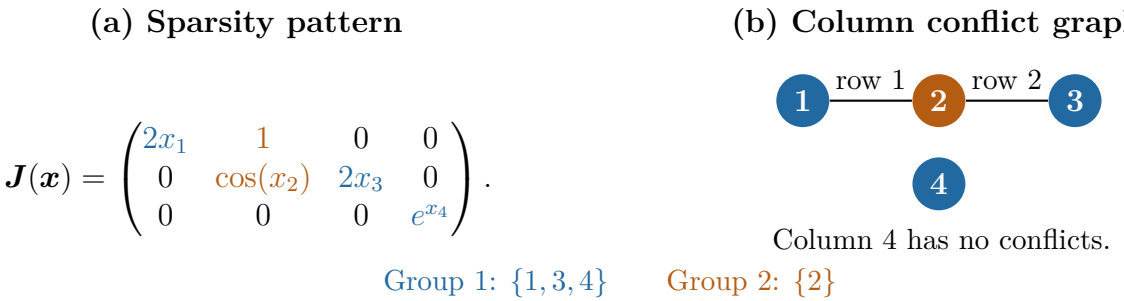


Figure 15: The same four-input example viewed as a pattern and as a graph. Colored coefficients mark possible nonzero entries. Graph edges record conflicts between columns, with the responsible row written beside each edge. Blue columns 1, 3, and 4 have no mutual conflicts and share one seed; orange column 2 uses another. A color denotes a group of columns.

the graph. The number of colors in the graph corresponds to the number of JVP evaluations needed to recover all columns of $\boldsymbol{J}$. The columns of the same color have no conflicts and can be computed together in a single JVP evaluation. In our example, $\chi(G) = 2$, as shown in Fig. 15. Columns 1, 3, and 4 can be assigned the same color (blue), while column 2 can be assigned another color (orange). As we saw previously, we could recover columns 1, 3, and 4 with a single JVP evaluation along $\boldsymbol{e}_1 + \boldsymbol{e}_3 + \boldsymbol{e}_4$, and column 2 with another JVP evaluation along $\boldsymbol{e}_2$. Which is consistent with the coloring we found. For this simple example, the chromatic number is easy to find by inspection. But in general, finding the chromatic number of a graph is an NP-hard problem. Therefore, we need to use certain algorithms to find a coloring that uses a small number of colors, even if it is not the minimum possible.

B.4 Decompression

After finding a coloring, we can evaluate the JVPs along the vectors belonging to each vector. We can group them in a matrix $\boldsymbol{S} \in \mathbb{R}^{m \times c}$ called the *seed matrix*. For our example, the seed matrix is

$$\boldsymbol{S} = \begin{pmatrix} 1 & 0 \\ 0 & 1 \\ 1 & 0 \\ 1 & 0 \end{pmatrix}, \tag{B.9}$$

We can then get a compressed Hessian by computing the product $\boldsymbol{J}(\boldsymbol{x})\boldsymbol{S}$. For our example, we get

$$\boldsymbol{J}(\boldsymbol{x})\boldsymbol{S} = \begin{pmatrix} 2x_1 & 1 \\ 2x_3 & \cos(x_2) \\ e^{x_4} & 0 \end{pmatrix}. \tag{B.10}$$

This compressed matrix contains all the information needed to recover the original Jacobian. Instead of storing a full 3×4 Jacobian, we only need to store a 3×2 compressed matrix. To recover the original $\boldsymbol{J}$, we can use the coloring information to unpack the compressed matrix. The first compressed column supplies the blue entries, and the second supplies the orange entries:

$$
\begin{aligned}
J_{11} = Y_{11} &= 2x_1, & J_{12} = Y_{12} &= 1, \\
J_{23} = Y_{21} &= 2x_3, & J_{22} = Y_{22} &= \cos(x_2), \\
J_{34} = Y_{31} &= e^{x_4}.
\end{aligned} \tag{B.11}
$$

The sparsity pattern places these values in their original columns and sets all other entries to zero:

$$
\boldsymbol{J}(\boldsymbol{x}) = \begin{pmatrix} 2x_1 & 1 & 0 & 0 \\ 0 & \cos(x_2) & 2x_3 & 0 \\ 0 & 0 & 0 & e^{x_4} \end{pmatrix}. \tag{B.12}
$$

Recovering $\boldsymbol{J}$ from $\boldsymbol{J}(\boldsymbol{x})\boldsymbol{S}$ is called *decompression*.

B.5 The Hessian and HVPs

For a function $f : \mathbb{R}^m \to \mathbb{R}$, the Hessian matrix is defined as

$$
\boldsymbol{H}(\boldsymbol{x}) = \begin{pmatrix} \frac{\partial^2 f}{\partial x_1^2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_m} \\ \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_m \partial x_1} & \cdots & \frac{\partial^2 f}{\partial x_m^2} \end{pmatrix} \in \mathbb{R}^{m \times m}. \tag{B.13}
$$

The methods we have discussed for the Jacobian can also be applied to the Hessian. In fact, it can be applied to any sparse matrix, regardless of whether it is a Jacobian or Hessian. Therefore, the same coloring and decompression techniques can be used to recover the Hessian with fewer evaluations. Since the Hessian is symmetric, we can exploit this property to further reduce the number of evaluations needed. Furthermore, computing the Hessian-vector product (HVP) is equivalent to computing the vector-Hessian product (VHP) due to the symmetry of the Hessian. For instance, suppose we have a function $f : \mathbb{R}^4 \to \mathbb{R}$ defined as

$$
f(\boldsymbol{x}) = x_1^2 + x_1 x_2 + x_3^2 + x_4^2, \tag{B.14}
$$

Its Hessian is given by

$$
\boldsymbol{H}(\boldsymbol{x}) = \begin{pmatrix} 2 & 1 & 0 & 0 \\ 1 & 0 & 0 & 0 \\ 0 & 0 & 2 & 0 \\ 0 & 0 & 0 & 2 \end{pmatrix}. \tag{B.15}
$$

Using the same graph coloring technique as before, we can find a coloring for the Hessian sparsity pattern. Columns 1 and 2 are the only ones that have a nonzero entry in the same row, so they cannot be assigned the same color. So $\chi(G) = 2$ and we can again assign blue to columns 1, 3, and 4 and orange to column 2. Colors follow columns, so the symmetric entries H_{12} and H_{21} have different colors despite having the same value. Therefore, we can get the seed matrix $\boldsymbol{S}$ as

$$\boldsymbol{S} = \begin{pmatrix} 1 & 0 \\ 0 & 1 \\ 1 & 0 \\ 1 & 0 \end{pmatrix}. \tag{B.16}$$

Then we can get the compressed Hessian by computing the product $\boldsymbol{H}(\boldsymbol{x})\boldsymbol{S}$

$$\boldsymbol{H}(\boldsymbol{x})\boldsymbol{S} = \begin{pmatrix} 2 & 1 \\ 1 & 0 \\ 2 & 0 \\ 2 & 0 \end{pmatrix}. \tag{B.17}$$

To recover the original Hessian, we can use the coloring information to unpack the compressed matrix

$$\begin{aligned} H_{11} = Y_{11} = 2, \qquad & H_{12} = Y_{12} = 1, \\ H_{21} = Y_{21} = 1, \qquad & H_{33} = Y_{31} = 2, \\ H_{44} = Y_{41} = 2. & \end{aligned} \tag{B.18}$$

All other coefficients are structural zeros.

B.6 ASD in PET

Since the potential energy computed by PET is the sum of the contributions from each atomic environment A_i, an atom i can only interact with its neighbors in A_i. Therefore, when computing the second derivatives of the potential, only atoms inside A_i can contribute to the Hessian entries $H_{i\alpha,j\beta}$, the other entries are structural zeros. The Hessian can be written as blocks of 3×3 matrices, where each block corresponds to the second derivatives of the potential with respect to the coordinates of two atoms i and j:

$$H_{i\alpha,j\beta} = \sum_t \frac{\partial^2 \varepsilon_t}{\partial R_{i\alpha} \partial R_{j\beta}}. \tag{B.19}$$

We can represent the pairs of atoms that interact with each other as a graph, where each atom is a vertex and an edge connects two atoms if they are neighbors in the same atomic environment. A compact way of visualizing these interactions is by using an adjacency matrix $\boldsymbol{A}$, where $A_{ij} = 1$ if atoms i and j are connected, and 0 otherwise. It is important to construct this matrix, because by doing so we can identify the nonzero entries of the Hessian and thus use the graph coloring technique to detect the sparsity pattern of the Hessian, without having to compute the Hessian explicitly. With this graph, we can identify paths between different atoms, even if they are not directly connected. This is achieved by traversing intermediate nodes, which allows us to find all pairs of atoms that are connected by a path of length k or less. The number of nodes to get from one atom to another is called a hop. We can determine the k-hop pattern by computing the powers of the adjacency matrix $\boldsymbol{A}$.

$$\boldsymbol{A}^{(k)} = \bigvee_{i=1}^{k} \boldsymbol{A}^i \tag{B.20}$$

The coefficients $\boldsymbol{A}_{ij}^{(k)}$ are 1 if atoms i and j are connected by a path of length k or less, and 0 otherwise. For instance, if $(\boldsymbol{A}^{(2)})_{ij} = 1$, atoms i and j are connected by a path of length 2 or less.

B.7 Hessian truncation

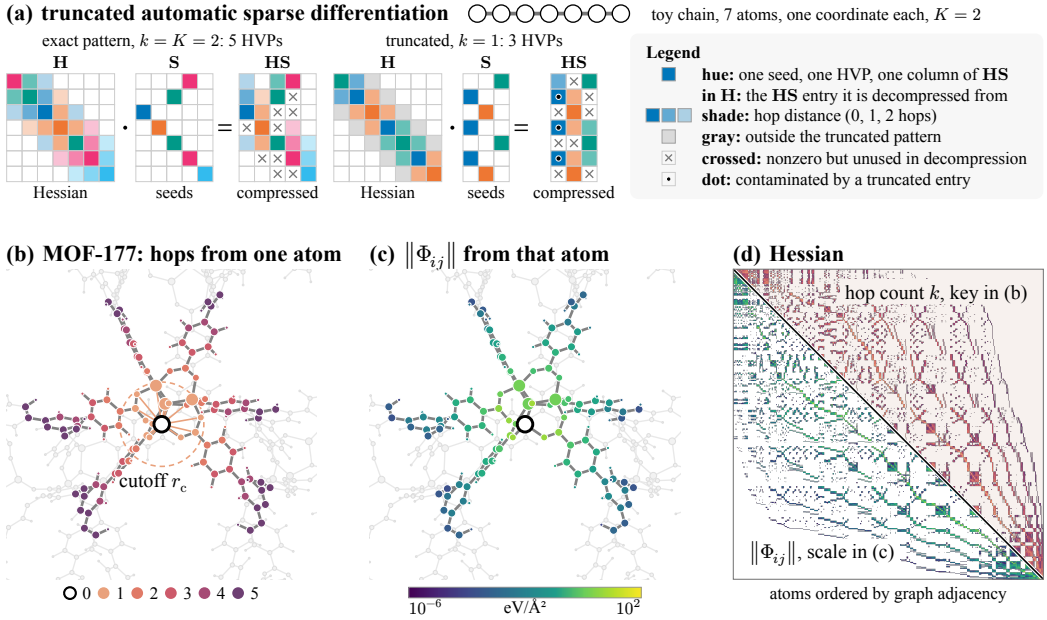


Figure 16: Truncation error in the Hessian. The blue blocks are retained, and the orange blocks are discarded. The retained blocks are contaminated by the omitted couplings, which are still present in the energy function. Figure adapted from [43].

The number of hops K needed to recover the full Hessian is given by $K = 2L + 1$, where L is the number message-passing layers in the PET architecture. For instance, PET-MAD-1.5-S has $L = 3$ message-passing layers, so we need $K = 7$ hops to recover the full Hessian. Therefore, $\boldsymbol{A}^{(7)}$ tell us which entries are nonzero and we can use it to find the coloring of the Hessian. The higher K is, the more nonzero entries we have in the Hessian, and the more evaluations we need to recover it. If we truncate the Hessian to a smaller number of hops $k < K$, $\boldsymbol{A}^{(k)}$ becomes sparser and fewer HVPs are needed to recover the truncated Hessian. It has been shown in [43] that this is a reasonable approximation, since the Hessian entries decay with distance and it allows for a more efficient computation. We can see this in Fig. 16 on plot (b), where we have the atoms around the central atom, each having a color corresponding to the number of hops from the central atom. We see that all the one hop neighbors are located inside the cutoff radius. The further away we are from the central atom, the higher the number of hops. In the plot (c), we have the frobenius norm of the Hessian blocks for the previous example between the central atom and its neighbors. It is clear from this plot that the norm decays with the number of hops, thus justifying the truncation of the Hessian.

C Temperature and pressure control

C.1 Thermostat

A thermostat exchanges energy with the particles to sample a target temperature T. The equipartition theorem states that each independent quadratic degree of freedom contributes $k_{\mathrm{B}}T/2$ to the mean energy [6]:

$$\left\langle \frac{ax^2}{2} \right\rangle = \frac{1}{2}k_{\mathrm{B}}T, \qquad a > 0. \tag{C.1}$$

For f active atomic momentum degrees of freedom, this gives

$$K = \sum_{i=1}^{N} \frac{|\boldsymbol{P}_i|^2}{2M_i}, \qquad \langle K \rangle = \frac{f}{2}k_{\mathrm{B}}T. \tag{C.2}$$

Thus

$$T = \frac{2\langle K \rangle}{f k_{\mathrm{B}}}, \qquad T_{\mathrm{kin}} = \frac{2K}{f k_{\mathrm{B}}}. \tag{C.3}$$

Here $f = 3N$ for unconstrained momenta, or $3N - 3$ when total linear momentum is removed. The instantaneous T_{kin} fluctuates around the target.

Nosé introduces a positive generalized coordinate s and its conjugate momentum p_s [36]. At fixed cell, the extended Hamiltonian is

$$\mathcal{H}_{\mathrm{N}}(\boldsymbol{R}^N, s; \boldsymbol{\pi}^N, p_s) = \sum_{i=1}^{N} \frac{|\boldsymbol{\pi}_i|^2}{2M_i s^2} + U(\boldsymbol{R}^N) + \frac{p_s^2}{2Q} + f k_{\mathrm{B}}T \ln s. \tag{C.4}$$

The generalized coordinates are $(\boldsymbol{R}^N, s)$ and the canonical momenta are $(\boldsymbol{\pi}^N, p_s)$; Q is the thermostat inertia. The coefficient f is chosen for canonical sampling in physical time. Hamilton's equations in Nosé's auxiliary time τ are

$$\frac{\mathrm{d}\boldsymbol{R}_i}{\mathrm{d}\tau} = \frac{\partial \mathcal{H}_{\mathrm{N}}}{\partial \boldsymbol{\pi}_i} = \frac{\boldsymbol{\pi}_i}{M_i s^2}, \tag{C.5}$$

$$\frac{\mathrm{d}\boldsymbol{\pi}_i}{\mathrm{d}\tau} = -\frac{\partial \mathcal{H}_{\mathrm{N}}}{\partial \boldsymbol{R}_i} = \boldsymbol{F}_i, \qquad \boldsymbol{F}_i = -\nabla_{\boldsymbol{R}_i} U, \tag{C.6}$$

$$\frac{\mathrm{d}s}{\mathrm{d}\tau} = \frac{\partial \mathcal{H}_{\mathrm{N}}}{\partial p_s} = \frac{p_s}{Q}, \tag{C.7}$$

$$\frac{\mathrm{d}p_s}{\mathrm{d}\tau} = -\frac{\partial \mathcal{H}_{\mathrm{N}}}{\partial s} = \sum_i \frac{|\boldsymbol{\pi}_i|^2}{M_i s^3} - \frac{f k_{\mathrm{B}}T}{s}. \tag{C.8}$$

Using physical time $\mathrm{d}t = \mathrm{d}\tau/s$, physical momenta $\boldsymbol{P}_i = \boldsymbol{\pi}_i/s$, and $\eta = \ln s$ gives the Nosé–Hoover equations [37]:

$$\dot{\boldsymbol{R}}_i = \frac{\boldsymbol{P}_i}{M_i}, \qquad \dot{\boldsymbol{P}}_i = \boldsymbol{F}_i - \dot{\eta}\boldsymbol{P}_i, \qquad Q\ddot{\eta} = 2K - f k_{\mathrm{B}}T. \tag{C.9}$$

Dots now denote physical-time derivatives. The thermostat removes excess kinetic energy by increasing the friction $\dot{\eta}$, and reverses this feedback when the kinetic energy is below its target.

The project uses a chain of three thermostat variables to improve sampling [38]. Introduce coordinates $\boldsymbol{\eta} = (\eta_1, \eta_2, \eta_3)$ and auxiliary momenta $p_{\eta_j} = Q_j \dot{\eta}_j$. Its conserved extended energy, often called the chain Hamiltonian, is

$$\begin{aligned} &\mathcal{H}_{\mathrm{NHC}}(\boldsymbol{R}^N, \eta_1, \eta_2, \eta_3; \boldsymbol{P}^N, p_{\eta_1}, p_{\eta_2}, p_{\eta_3}) \\ &\quad = K + U(\boldsymbol{R}^N) + \sum_{j=1}^{3} \frac{p_{\eta_j}^2}{2Q_j} + k_{\mathrm{B}}T(f\eta_1 + \eta_2 + \eta_3). \end{aligned} \tag{C.10}$$

Keeping only the first thermostat recovers the single Nosé–Hoover extended energy. In these physical-time variables the dynamics is noncanonical: ordinary Hamilton's equations applied to this energy alone would omit the friction couplings. The corresponding generalized equations are

$$\dot{\boldsymbol{R}}_i = \frac{\boldsymbol{P}_i}{M_i}, \qquad \dot{\boldsymbol{P}}_i = \boldsymbol{F}_i - \frac{p_{\eta_1}}{Q_1}\boldsymbol{P}_i, \tag{C.11}$$

$$\dot{\eta}_j = \frac{p_{\eta_j}}{Q_j}, \qquad j = 1, 2, 3, \tag{C.12}$$

$$\dot{p}_{\eta_1} = 2K - fk_{\mathrm{B}}T - \frac{p_{\eta_2}}{Q_2}p_{\eta_1}, \tag{C.13}$$

$$\dot{p}_{\eta_2} = \frac{p_{\eta_1}^2}{Q_1} - k_{\mathrm{B}}T - \frac{p_{\eta_3}}{Q_3}p_{\eta_2}, \tag{C.14}$$

$$\dot{p}_{\eta_3} = \frac{p_{\eta_2}^2}{Q_2} - k_{\mathrm{B}}T. \tag{C.15}$$

Each additional thermostat controls the kinetic energy of the preceding one. Equivalently, $\xi_j = \dot{\eta}_j = p_{\eta_j}/Q_j$ are the friction variables.

C.2 Barostat

A barostat on the other hand exchanges work with the simulation cell to keep the system at a target pressure P_{ext}. The barostat adjusts the simulation cell size and shape to maintain the desired pressure. This project uses the flexible-cell formulation of Shinoda, Shiga, and Mikami [35], combining Martyna–Tobias–Klein (MTK) pressure control [34] with Parrinello–Rahman strain dynamics, as implemented in LAMMPS `fix npt` [14]. The instantaneous pressure includes kinetic and configurational terms:

$$\boldsymbol{\Pi}_{\mathrm{int}} = \frac{1}{V}\left[\sum_i \frac{\boldsymbol{P}_i \otimes \boldsymbol{P}_i}{m_i} + \boldsymbol{W}_{\mathrm{vir}}\right], \qquad P_{\mathrm{int}} = \frac{1}{3}\operatorname{tr}\boldsymbol{\Pi}_{\mathrm{int}}, \tag{C.16}$$

where $\boldsymbol{W}_{\mathrm{vir}}$ is the configurational virial tensor. The feedback is illustrated by the isotropic MTK equations. Let $\nu = \dot{V}/(3V)$ be the linear expansion rate, W_b the fictitious barostat inertia, and ξ_b the first friction variable of the chain thermostatting the barostat. For an imposed pressure P_{ext},

$$\dot{\boldsymbol{R}}_i = \frac{\boldsymbol{P}_i}{m_i} + \nu \boldsymbol{R}_i, \qquad \dot{V} = 3\nu V, \tag{C.17}$$

$$\dot{\boldsymbol{P}}_i = \boldsymbol{F}_i - \left[\xi_1 + \left(1 + \frac{3}{f}\right)\nu\right] \boldsymbol{P}_i, \tag{C.18}$$

$$W_b \dot{\nu} = 3V(P_{\text{int}} - P_{\text{ext}}) + \frac{6K}{f} - \xi_b W_b \nu. \tag{C.19}$$

The pressure difference accelerates the cell outward or inward; cell motion rescales positions and momenta and exchanges mechanical work with the atoms. The $3/f$ momentum correction and $6K/f$ barostat term ensure the intended NPT statistical measure [34].